



# Trees

Fall 2020

School of Software Engineering  
South China University of Technology

# Contents

- Definitions of tree
- Binary tree
- AVL tree
- **Splay tree**
- B-tree

# Splay Trees

# Readings

- Reading
  - Sections 4.5-4.7

# Self adjusting Trees

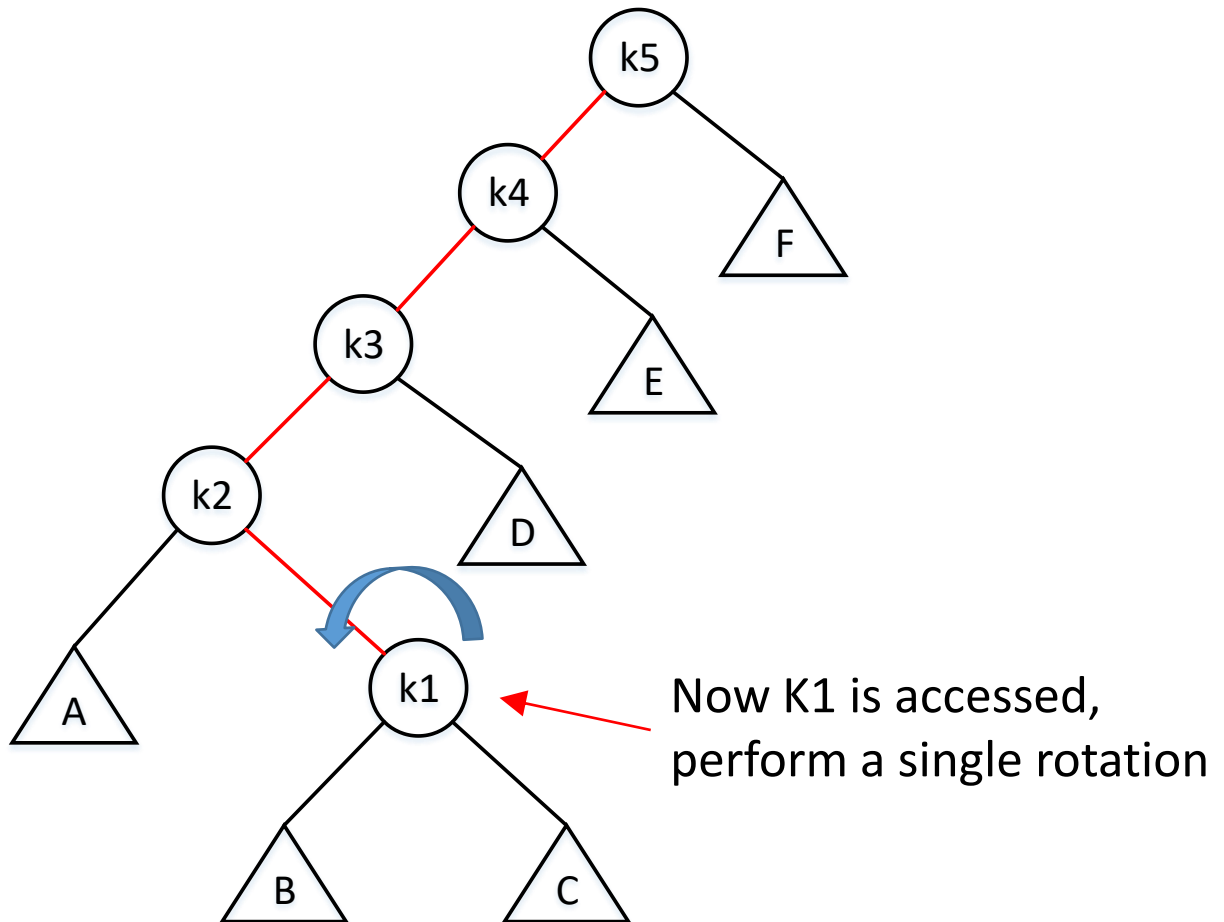
- Ordinary binary search trees have no balance conditions
  - what you get from insertion order is it
- Balanced trees like AVL trees enforce a balance condition when nodes change
  - tree is always balanced after an insert or delete
- Self-adjusting trees get reorganized over time as nodes are accessed
  - Tree adjusts after insert, delete, or find

# Splay Trees

- Splay trees are tree structures that:
  - Are not perfectly balanced all the time
  - Data most recently accessed is near the root. (principle of locality; 80-20 “rule”)
- The procedure:
  - After node X is accessed, perform “**splaying**” operations to bring X to the root of the tree.
  - Do this in a way that leaves the tree more balanced as a whole


# Splay Trees

- How to perform splaying?
  - A Simple Idea
  - to perform single rotations, bottom up



# Splay Trees

- How to perform splaying?
  - A Simple Idea
  - to perform single rotations, bottom up




Then, continue to perform a single rotation



# Splay Trees


- How to perform splaying?
  - A Simple Idea
  - to perform single rotations, bottom up



Then, continue...

# Splay Trees


- How to perform splaying?
  - A Simple Idea
  - to perform single rotations, bottom up



continue ...

# Splay Trees

- How to perform splaying?
  - A Simple Idea
  - to perform single rotations, bottom up




Now k1 reach the root.  
Future accesses on k1  
are cheaper.

But pushing other node  
deep, ex. k3.

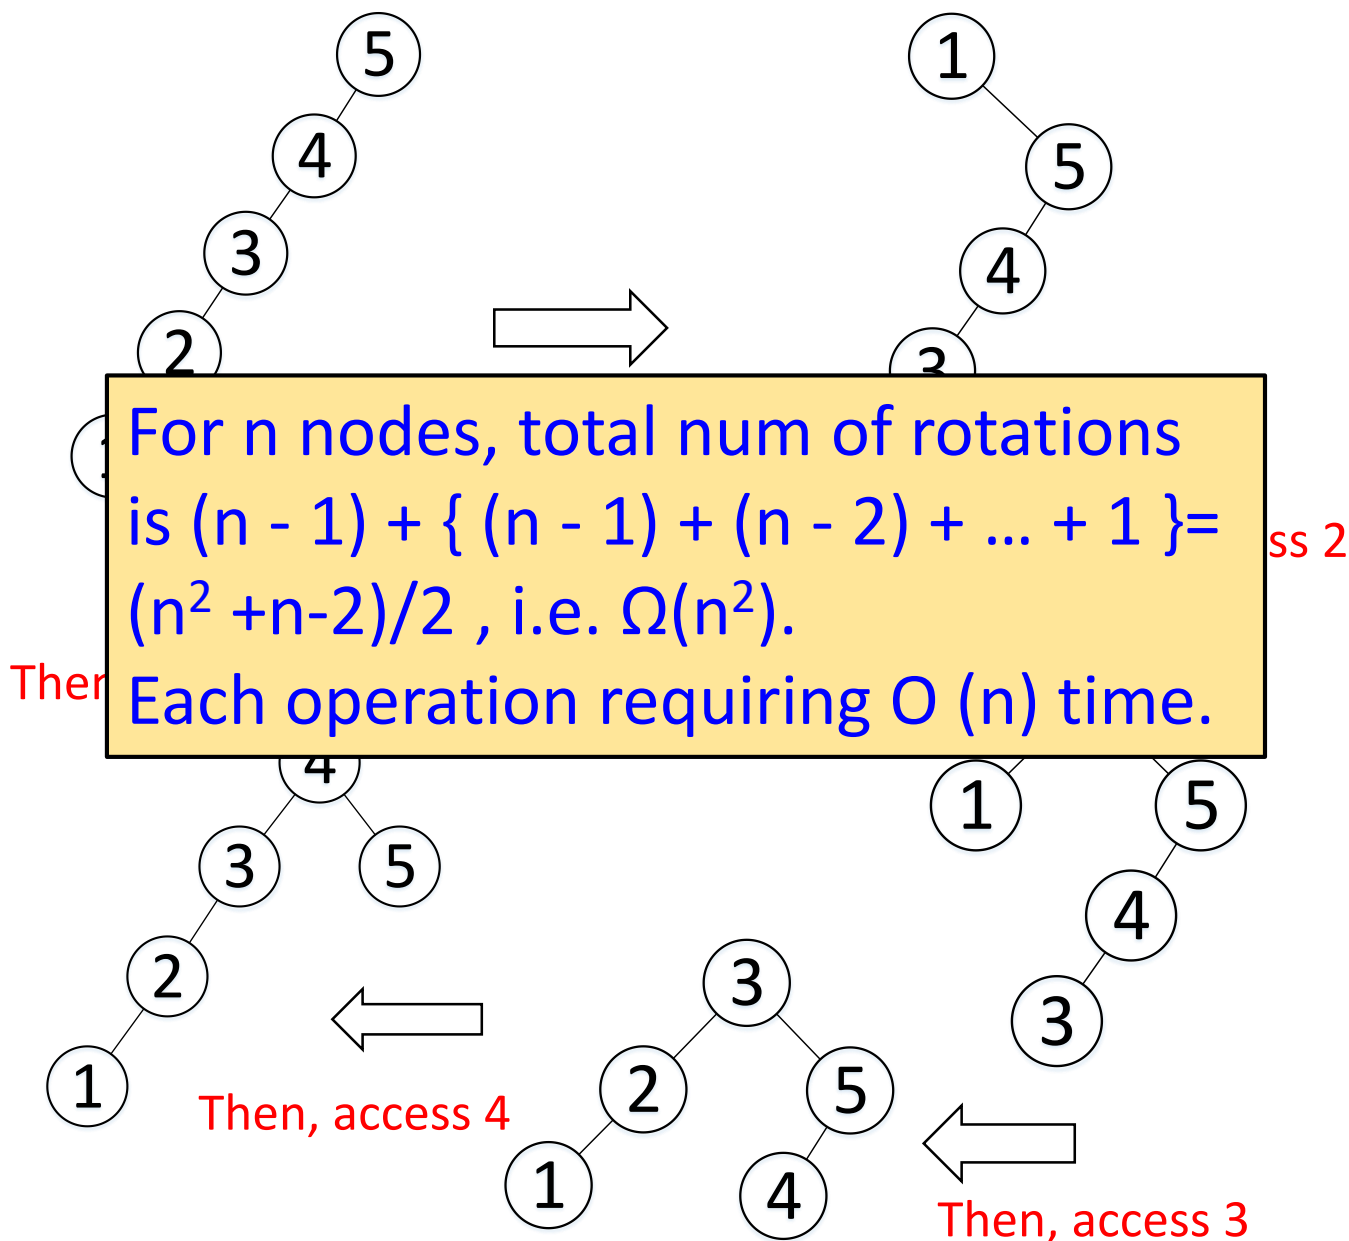
# Splay Trees

- How to perform splaying?
  - A Simple Idea
  - to perform single rotations, bottom up
- This strategy is not good enough
  - A sequence of  $M$  operations requires  $\Omega(N \cdot M)$



Consider the tree formed by inserting keys 1, 2, 3, ...,  $N$  into an initially empty tree

# Splay Trees




# Splay Trees

- Goal:
  - Principle of locality
  - An  $O(\log N)$  amortized cost per operation
- How to perform splaying?
  - Another strategy to perform double rotations (splay)

# Splay Trees


- Terminology

- Let  $X$  be a non-root node with  $\geq 2$  ancestors.
  - $P$  is its parent node.
  - $G$  is its grandparent node.




# Zig-Zig and Zig-Zag

Parent and grandparent in same direction.



Parent and grandparent in different directions.






# Splay Trees

## Splay Tree Operations:

1. Helpful if nodes contain a **parent** pointer.




2. When X is accessed, apply one of **six** rotation routines.

- Single Rotations (X has a P (the root) but no G)  
    ZigFromLeft, ZigFromRight
- Double Rotations (X has both a P and a G)  
    ZigZigFromLeft, ZigZigFromRight  
    ZigZagFromLeft, ZigZagFromRight

# Zig at depth 1 (single rotation)


- “Zig” is just a **single rotation**, as in an AVL tree
- Let R be the node that was accessed (e.g. using Find)



- ZigFromLeft moves R to the top → faster access next time

# Zig at depth 1 (single rotation)


- Suppose Q is now accessed using Find



- ZigFromRight moves Q back to the top

# Zig-Zag operation

- “Zig-Zag” consists of **two rotations of the opposite direction** (assume R is the node that was accessed)



(ZigFromRight)


(ZigFromLeft)

ZigZagFromLeft




# Zig-Zig operation

- “Zig-Zig” consists of **two single rotations of the same direction** (R is the node that was accessed)



# Decreasing depth - "autobalance"



Find(T)

Find(R)


- splaying not only moves the accessed node to the root,
- but also has the effect of **roughly halving** the depth of most nodes on the access path

# Splay Tree Insert and Delete

- Insert x
  - Insert x as normal then splay x to root.
- Delete x
  - Splay x to root and remove it. (note: the node does not have to be a leaf or single child node like in BST delete.) Two trees remain, right subtree and left subtree.
  - Splay the max in the left subtree to the root
  - Attach the right subtree to the new root of the left subtree.

# Example Insert

- Inserting in order 1,2,3,...,8
- Without self-adjustment

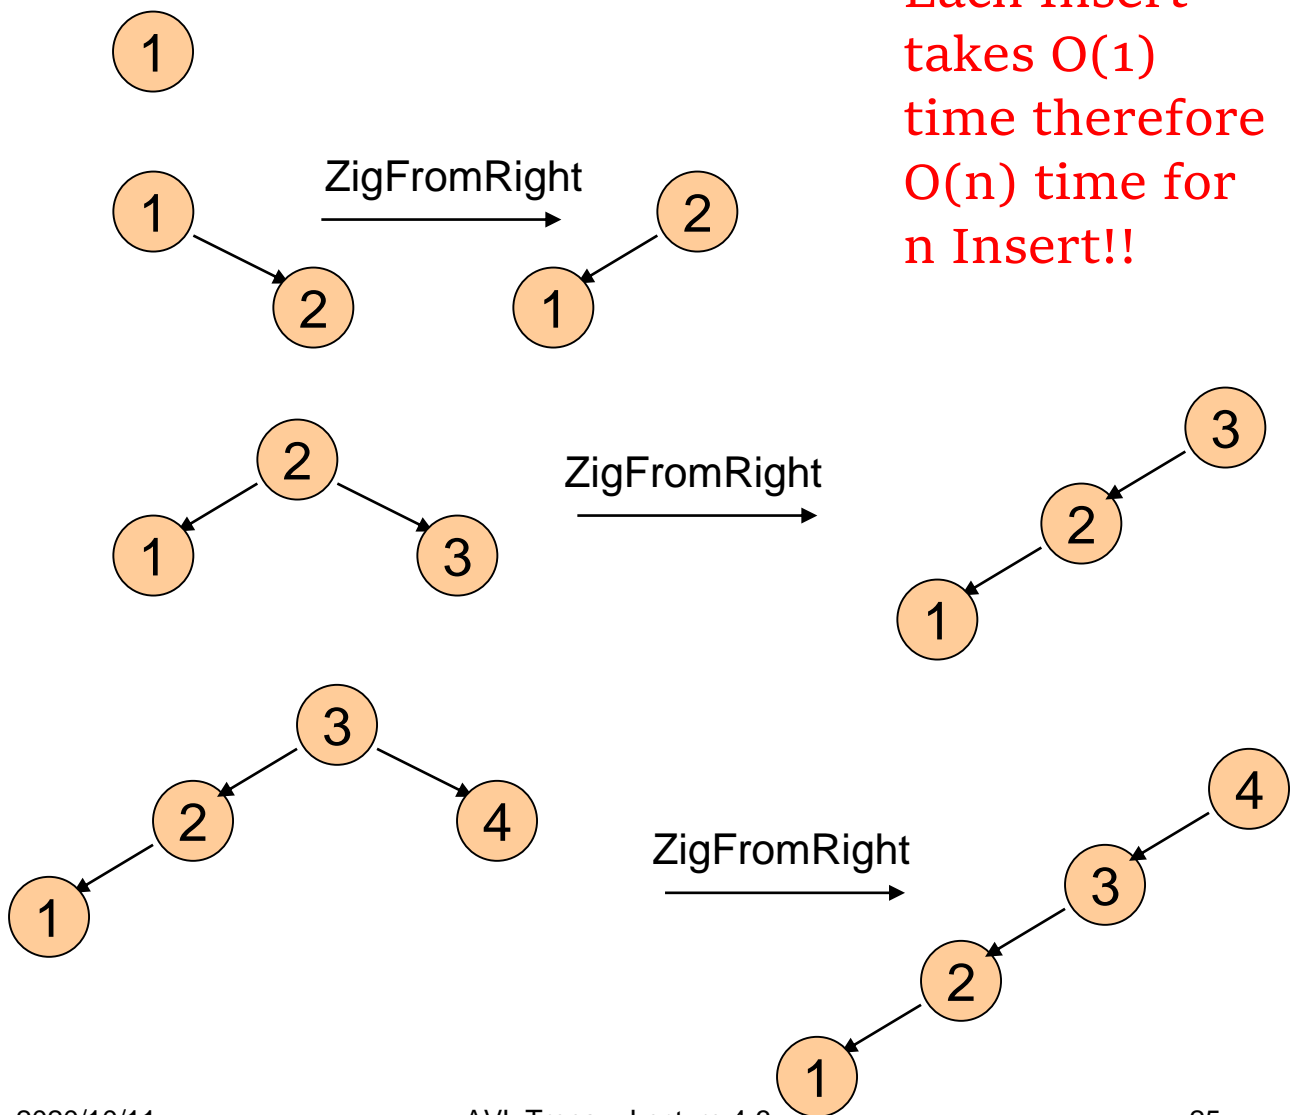





# Example Insert

- Inserting in order 1,2,3,...,8
- With self-adjustment

Each Insert takes  $O(1)$  time therefore  $O(n)$  time for  $n$  Insert!!



# Example Deletion



# Analysis of Splay Trees

- Splay trees tend to be balanced
  - $M$  operations takes time  $O(M \log N)$  for  $M \geq N$  operations on  $N$  items. (proof is difficult)
  - Amortized  $O(\log n)$  time.
- Splay trees have good “locality” properties
  - Recently accessed items are near the root of the tree.
  - Items near an accessed one are pulled toward the root.


# Homework

- Homework 3-3
  - Textbook exercises 4.27, 4.28

# B-Trees

# Beyond Binary Search Trees: Multi-Way Trees

- Example: B-tree of order 3 has 2 or 3 children per node
- e.g. search for 8



# B-Trees

- B-Trees are multi-way search trees commonly used in database systems or other applications where data is stored **externally on disks** and **keeping the tree shallow** is important.

# B-Trees

- A B-Tree of order  $M$  has the following properties:
  - The **root** is either a leaf or has **between 2 and  $M$  children**.
  - All **nonleaf nodes** (except the root) have **between  $\lceil M/2 \rceil$  and  $M$  children**.
  - All leaves are at the same depth.


- All data records are stored at the leaves.
- Internal nodes have “keys” guiding to the leaves.
- Leaves store between  $\lceil L/2 \rceil$  and  $L$  data records, where  $L$  can be equal to  $M$  (default) or can be different.



# B-Tree Details

Each (non-leaf) internal node of a B-tree has:


- Between  $\lceil M/2 \rceil$  and  $M$  children.
- up to  $M-1$  **keys**  $k_1 < k_2 < \dots < k_{M-1}$



Keys are ordered so that:

$$k_1 < k_2 < \dots < k_{M-1}$$

# Properties of B-Trees




- Children of each internal node are "between" the items in that node.

Suppose subtree  $T_i$  is the  $i$ th child of the node:


- All keys in  $T_i$  must be between keys  $k_{i-1}$  and  $k_i$   
i.e.  $k_{i-1} \leq T_i < k_i$ ,  $k_{i-1}$  is the smallest key in  $T_i$
- All keys in first subtree  $T_1 < k_1$
- All keys in last subtree  $T_M \geq k_{M-1}$

# Properties of B-Trees

## B-Tree Nonleaf Node




- The  $K$ s are keys
- The  $P$ s are pointers to subtrees

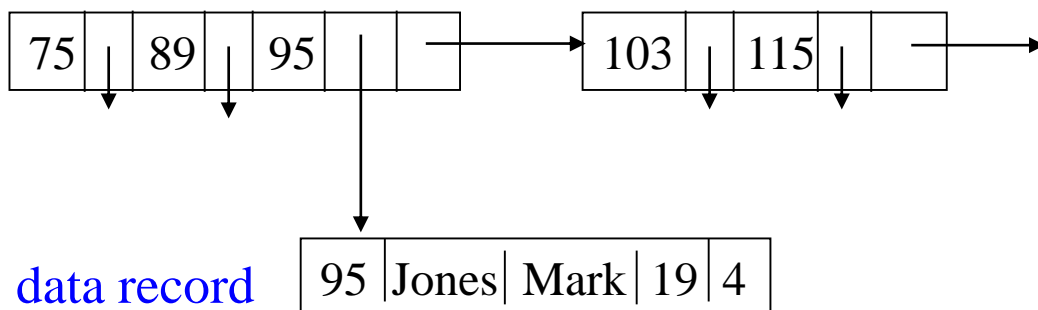


# Properties of B-Trees

## B-Tree leaf Node (B+tree)




- The Ks are keys (assume unique).
- The Rs are pointers to records with those keys.
- The Next link points to the next leaf in key order (B+-tree).



# Searching in B-trees

- B-tree of order 3: also known as 2-3 tree (2 to 3 children)



- Examples: Search for 9, 14, 12
- Note: If leaf nodes are connected as a Linked List, B-tree is called a B+ tree – Allows sorted list to be accessed easily

# Searching in B-trees

- Searching a B-Tree T for a Key Value K

```
Find(ElementType K, Btree T){
    B = T;
    while (B is not a leaf){
        find the Pi in node B that points to the proper
        subtree that K will be in;

        B = Pi;
    }


    /* Now we're at a leaf */

    if key K is the jth key in leaf B,
        use the jth record pointer to find the
        associated record;
    else /* K is not in leaf B */ report failure;
}
```

# Inserting into B-Trees

- Insert X: Do a Find on X and find appropriate leaf node
  - If leaf node is not full, fill in empty slot with X
    - E.g. Insert 5
  - If leaf node is full, **split** leaf node and adjust parents up to root node
    - E.g. Insert 9

Assume  $M=L=3$ ,  
so (6 7 8) is full.




# Inserting into B-Trees

// Inserting a New Key in a B-Tree of Order M (and L=M)

Insert(ElementType K, Btree B)

```
{  
    find the leaf node LB of B in which K belongs;  
    if notfull(LB) insert K into LB;  
    else {  
        split LB into two nodes LB and LB2 with  
         $j = \lfloor (M+1)/2 \rfloor$  keys in LB and the rest in LB2;
```



```
    if ( IsNull(Parent(LB)) )  
        CreateNewRoot(LB,  $K[j+1]$ , LB2);  
    else  
        InsertInternal(Parent(LB),  $K[j+1]$ , LB2);  
    }  
}
```



# Inserting into B-Trees

// Inserting a (Key,Ptr) Pair into an Internal Node

If the node is not full, insert them in the proper place and return.

If the node is already full (M pointers, M-1 keys), find the place for the new pair and split the adjusted (Key,Ptr) sequence into two internal nodes with

$j = \lfloor (M+1)/2 \rfloor$  pointers and  $j-1$  keys in the first,

the next key is inserted in the node's parent, and the rest in the second of the new pair.

# Example of Insertions


Insertions into a B+tree with  $M=3$ ,  $L=2$

Insertion Sequence: 9, 5, 1, 7, 3, 12


1




2




3




4



5




6



# Deleting From B-Trees

- Delete X : Do a find and remove from leaf
  - Leaf underflows – borrow from a neighbor
    - E.g. 11
  - Leaf underflows and can't borrow – merge nodes, delete parent
    - E.g. 17



# Run Time Analysis of B-Tree Operations

- For a B-Tree of order  $M$ 
  - Each internal node has up to  $M-1$  keys to search
  - Each internal node has between  $\lceil M/2 \rceil$  and  $M$  children
  - Depth of B-Tree storing  $N$  items is  $O(\log_{\lceil M/2 \rceil} N)$
- Find: Run time is:
  - $O(\log M)$  to binary search which branch to take at each node. But  $M$  is small compared to  $N$ .
  - Total time to find an item is  $O(\text{depth} * \log M) = O(\log N)$

# Run Time Analysis of B-Tree Operations

- How Do We Select the Order  $M$ ?
  - In internal memory, small orders, like 3 or 4 are fine.
  - On disk, we have to worry about the number of disk accesses to search the index and get to the proper leaf.

**Rule: Choose the largest  $M$  so that an internal node can fit into one physical block of the disk.**

- This leads to typical  $M$ 's between 32 and 256
- And keeps the trees as shallow as possible.

# Summary of Search Trees

- Problem with Binary Search Trees
  - Must keep tree balanced to allow fast access to stored items
- **AVL trees:** Insert/Delete operations keep tree balanced
- **Splay trees:** Repeated Find operations produce balanced trees
- **Multi-way search trees** (e.g. B-Trees):
  - More than two children per node allows shallow trees; all leaves are at the same depth.
  - Keeping tree balanced at all times.
  - Excellent for indexes in database systems.

# Homework

- Homework 3-4

- Show the updated B+-Tree with order 4 that results from inserting the records U and R in order.
- Assume that the leaf nodes are capable of storing up to 3 records

