



Graph Algorithms

Fall 2020


School of Software Engineering
South China University of Technology

Shortest-Path Algorithms

Section 9.3

Recall Path cost ,Path length

- **Path cost**: the sum of the costs of each edge
- **Path length**: the number of edges in the path
 - Path length is the unweighted path cost




$$\text{length}(p) = 5$$

$$\text{cost}(p) = 11$$

Shortest Path Problems

- Given a graph $G = (V, E)$ and a “source” vertex s in V , find the minimum cost paths from s to every vertex in V
 - Single-Source Shortest Paths problem
- Many variations:
 - unweighted vs. weighted
 - cyclic vs. acyclic
 - pos. weights only vs. pos. and neg. weights
 - etc




Why study shortest path problems?

- **Traveling on a budget:** What is the cheapest airline schedule from Guangzhou to city X?
- **Optimizing routing of packets on the internet:**
 - Vertices are routers and edges are network links with different delays. What is the routing path with smallest total delay?
- **Shipping:** Find which highways and roads to take to minimize total delay due to traffic
- etc.

Unweighted Shortest Path


Problem: Given a “source” vertex s in an unweighted directed graph $G = (V, E)$, find the shortest path from s to all vertices in G

Only interested in path lengths



Breadth-First Search Solution

- **Basic Idea:** Starting at node s , find vertices that can be reached using 0, 1, 2, 3, ..., $N-1$ edges (works even for cyclic graphs!)



Breadth-First Search Algorithm

- Uses a queue to track vertices that are “nearby”
- source vertex is **s**
- Running time = $O(|V| + |E|)$

Distance[s] := 0

Enqueue(Q,s);

Mark(s); //After a vertex is marked once
// it won't be enqueued again

while queue is not empty do

 X := Dequeue(Q);

 for each vertex Y adjacent to X do

 if Y is unmarked then


 Distance[Y] := Distance[X] + 1;

 Previous[Y] := X; //if we want to record paths

 Enqueue(Q,Y);


 Mark(Y);

Example: Shortest Path length



Queue Q = C

Example (ct'd)




Queue $Q = A D E$

Indicates the vertex is marked 


Previous pointer 

Example (ct'd)




Q = D E B

Example (ct'd)




Q = B G

Example (ct'd)



Q = F


Example (ct'd)



Q = H

Weighted Shortest Path

- What if edges have weights?
- Breadth First Search does not work anymore
 - minimum *cost* path may have more edges than minimum *length* path



from C to A

Shortest path (length) = $C \rightarrow A$ (cost = 9)

Minimum Cost Path = $C \rightarrow E \rightarrow D \rightarrow A$ (cost = 8)

Dijkstra's Algorithm

- Classic algorithm for solving Single-Source Shortest Paths in weighted graphs (without negative weights)
- A **greedy algorithm** (irrevocably makes decisions without considering future consequences)


Basic Idea of Dijkstra's Algorithm

- Find the vertex with smallest cost that has not been “marked” yet.
- Mark it and compute the cost of its neighbors.
- Do this until all vertices are marked.
- Note that each step of the algorithm we are marking one vertex and we won't change our decision: hence the term “greedy” algorithm

Dijkstra's Shortest Path Algorithm

- Initialize the cost of s to 0, and all the rest of the nodes to ∞
- Initialize set S to be \emptyset
 - S is the set of nodes to which we have a shortest path
- While S is not all vertices
 - Select the node A with the lowest cost that is not in S and identify the node as now being in S
 - for each node B adjacent to A
 - if $\text{cost}(A) + \text{cost}(A, B) < B$'s currently known cost
 - set $\text{cost}(B) = \text{cost}(A) + \text{cost}(A, B)$
 - set $\text{previous}(B) = A$ so that we can remember the path


Example: Initialization



Pick vertex not in S with lowest cost in every step.


	S	1	2	3	4	5	6	7
initial	{}	<u>0</u>	∞	∞	∞	∞	∞	∞

Example: Update Cost neighbors




S	1	2	3	4	5	6	7
{}	<u>0</u>	∞	∞	∞	∞	∞	∞
{1}	<u>0</u>	2	∞	<u>1</u>	∞	∞	∞

Example




S	1	2	3	4	5	6	7
{}	<u>0</u>	∞	∞	∞	∞	∞	∞
{1}	<u>0</u>	2	∞	<u>1</u>	∞	∞	∞

Example




S	1	2	3	4	5	6	7
{}	<u>0</u>	∞	∞	∞	∞	∞	∞
{1}	<u>0</u>	2	∞	<u>1</u>	∞	∞	∞
{1,4}	<u>0</u>	<u>2</u>	3	<u>1</u>	3	9	5

Example




S	1	2	3	4	5	6	7
{}	<u>0</u>	∞	∞	∞	∞	∞	∞
{1}	<u>0</u>	2	∞	<u>1</u>	∞	∞	∞
{1,4}	<u>0</u>	<u>2</u>	3	<u>1</u>	3	9	5
{1,4,2}	<u>0</u>	<u>2</u>	3	<u>1</u>	3	9	5

Example




S	1	2	3	4	5	6	7
{}	<u>0</u>	∞	∞	∞	∞	∞	∞
{1}	<u>0</u>	2	∞	<u>1</u>	∞	∞	∞
{1,4}	<u>0</u>	<u>2</u>	3	<u>1</u>	3	9	5
{1,4,2}	<u>0</u>	<u>2</u>	3	<u>1</u>	3	9	5
{1,4,2,5}	<u>0</u>	<u>2</u>	3	<u>1</u>	3	9	5

Example




S	1	2	3	4	5	6	7
{ }	<u>0</u>	∞	∞	∞	∞	∞	∞
{1}	<u>0</u>	2	∞	<u>1</u>	∞	∞	∞
{1,4}	<u>0</u>	<u>2</u>	3	<u>1</u>	3	9	5
{1,4,2}	<u>0</u>	<u>2</u>	3	<u>1</u>	3	9	5
{1,4,2,5}	<u>0</u>	<u>2</u>	3	<u>1</u>	<u>3</u>	9	5
{1,4,2,5,3}	<u>0</u>	<u>2</u>	<u>3</u>	<u>1</u>	<u>3</u>	8	5

Example



S	1	2	3	4	5	6	7
{}	<u>0</u>	∞	∞	∞	∞	∞	∞
{1}	<u>0</u>	2	∞	<u>1</u>	∞	∞	∞
{1,4}	<u>0</u>	<u>2</u>	3	<u>1</u>	3	9	5
{1,4,2}	<u>0</u>	<u>2</u>	3	<u>1</u>	3	9	5
{1,4,2,5}	<u>0</u>	<u>2</u>	3	<u>1</u>	3	9	5
{1,4,2,5,3}	<u>0</u>	<u>2</u>	3	<u>1</u>	3	8	5
{1,4,2,5,3,7}	<u>0</u>	<u>2</u>	3	<u>1</u>	3	<u>6</u>	5

Example



S	1	2	3	4	5	6	7
{}	<u>0</u>	∞	∞	∞	∞	∞	∞
{1}	<u>0</u>	2	∞	<u>1</u>	∞	∞	∞
{1,4}	<u>0</u>	<u>2</u>	3	<u>1</u>	3	9	5
{1,4,2}	<u>0</u>	<u>2</u>	3	<u>1</u>	3	9	5
{1,4,2,5}	<u>0</u>	<u>2</u>	3	<u>1</u>	3	9	5
{1,4,2,5,3}	<u>0</u>	<u>2</u>	3	<u>1</u>	3	8	5
{1,4,2,5,3,7}	<u>0</u>	<u>2</u>	3	<u>1</u>	3	<u>6</u>	5
{1,4,2,5,3,7,6}	<u>0</u>	<u>2</u>	3	<u>1</u>	3	<u>6</u>	5

Dijkstra's Algorithm - Implementation

//Implementation of Dijkstra's algorithm

//compute shortest path dists from "s"

void **Dijkstra**(Graph* G, int* D, int s) {

int i, v, w;

for (int i=0; i<G->n(); i++) // Initialize

D[i] = INFINITY;

D[0] = 0;

for (i=0; i<G->n(); i++) { //process vertices

//find the unvisited vertex with min dist

v = **minVertex**(G, D);

if (D[v] == INFINITY) return; //v is unreachable

G->setMark(v, VISITED);

//update the distance of v's neighbors

for (w=G->first(v); w<G->n(); w = G->next(v,w))

if (D[w] > (D[v] + G->weight(v, w)))

D[w] = D[v] + G->weight(v, w);

}

}

Dijkstra's Algorithm - Implementation

- **minVertex** – find the unvisited vertex with minimum distance
- **Method 1:** scan through the list of $|V|$ vertices searching for the minimum value

```
int minVertex(Graph* G, int* D) {  
    int i, v = -1;
```

```
//initialize v to some unvisited vertex  
for (i=0; i<G->n(); i++)  
    if (G->getMark(i) == UNVISITED) {  
        v = i;  
        break; }
```

```
// Now find smallest D value  
for (i++; i<G->n(); i++)  
    if ((G->getMark(i) == UNVISITED) && (D[i] < D[v]))  
        v = i;  
return v;  
}
```

Dijkstra's Algorithm - Implementation

- Complexity analysis of **method 1**
 - **minVertex** executes $|V|$ times, and it scans the list of $|V|$ vertices each time.
 - The cost is $\Theta(|V|^2)$
 - The edges are processed $\Theta(|E|)$ times, and each visit to an edge may cause a constant-time update to the array D.
 - The cost is $\Theta(|E|)$
- In total, the cost is $\Theta(|V|^2 + |E|) = \Theta(|V|^2)$

Dijkstra's Algorithm – Implementation

- Method 2:

- Store unprocessed vertices in a **priority-queue** (implemented using a **min-heap**) ordered by **distance values**.
- The next-closest vertex can be found in the heap in $\Theta(\log|V|)$ time
- Every time $D(x)$ is updated,
 - Reordered x in the heap by **deleting and reinserting** it.
 - Or, add the new smaller distance value for x as **a new record** in the heap
 - The greater distance values found later will be **ignored** because the vertex will already be marked as **VISITED**

Dijkstra's Algorithm – Implementation

// Implementation using the priority queue

// Class for elements in the heap

Class DijkElem {

Public:

int vertex, **distance**;

DijkElem() {vertex = -1; distance = -1;}

DijkElem(int v, int d) {vertex = v; distance = d};
};


```

//Implementation of Dijkstra's algorithm
void Dijkstra(Graph* G, int* D, int s) {
    int i, v, w;      // v is current vertex
    DijkElem temp;  DijkElem E[G->n()]; // Heap array

    for (int i=0; i<G->n(); i++) // Initialize
        D[i] = INFINITY;
    D[0] = 0;

    // Initialize heap array
    temp.distance = 0; temp.vertex = s;
    E[0] = temp;
    heap<DijkElem, DDComp> H(E, 1, G->n());

    // get an unvisited vertex with smallest distance
    for (i=0; i<G->n(); i++) {
        do {
            if(H.size() == 0) return; // Nothing to remove
            temp = H.removefirst(); //delmin
            v = temp.vertex;
        } while (G->getMark(v) == VISITED);

        G->setMark(v, VISITED); //mark the vertex

        if (D[v] == INFINITY) return; //unreachable

        for(w=G->first(v); w<G->n(); w=G->next(v,w))
            if (D[w] > (D[v] + G->weight(v, w))) {
                //update D
                D[w] = D[v] + G->weight(v, w);
                temp.distance = D[w]; temp.vertex = w;
                // Insert new distance in heap
                H.insert(temp);
            }
    }
}

```


Dijkstra's Algorithm – Implementation

- Complexity analysis of **method 2**
 - If the vertex with a smaller distance is inserted as a new record, the cost for finding the minimum value using the min-heap becomes $\Theta(\log |E|)$
 - It executes V times, and the cost is $\Theta(|V|\log |E|)$
 - Every visit to an edge cause an update to D and also the min-heap.
 - The cost is $\Theta(|E|\log |E|)$
- In total, the cost $\Theta((|V| + |E|)\log |E|)$

Correctness

- Dijkstra's algorithm is an example of a greedy algorithm
- Greedy algorithms always make choices that currently seem the best
 - Short-sighted – no consideration of long-term or global issues
 - Locally optimal does not always mean globally optimal
- In Dijkstra's case – choose the least cost node, but what if there is another path through other vertices that is cheaper?

“Cloudy” Proof: The Idea




- If the path to **G** is the next shortest path, the path to **P** must be at least as long. Therefore, any path through **P** to **G** cannot be shorter!

Inside the Cloud (Proof)

- Everything inside the cloud has the correct shortest path
- Proof is by **induction** on the number of nodes in the cloud:
 - **Base case:** Initial cloud is just the source s with shortest path 0.
 - **Inductive hypothesis:** Assume that a cloud of $k-1$ nodes all have shortest paths.
 - **Inductive step:** choose the least cost node G → has to be the shortest path to G (previous slide). Add k -th node G to the cloud.

Graphs with Negative Edge Costs

- If the graph has negative edge costs, then Dijkstra's algorithm does not work.



S	1	2	3	4	5
{}	<u>0</u>	∞	∞	∞	∞
{1}	<u>0</u>	210	∞	100	<u>20</u>
{1,5}	<u>0</u>	210	∞	<u>100</u>	20
{1,5,4}	<u>0</u>	210	<u>70</u>	100	20
{1,5,4,3}	<u>0</u>	<u>210</u>	70	100	20
{1,5,4,3,2}	<u>0</u>	<u>210</u>	70	60	20

All Pairs Shortest Path

- Given a edge weighted directed graph $G = (V, E)$, find for all u, v in V the length of the shortest path from u to v .
- Could run the appropriate single-source algorithm $|V|$ times.
 - On sparse graphs, it is fast to run $|V|$ Dijkstra's algorithms coded with priority queues
 - On dense graphs, the algorithm in textbook section 10.3.4 more faster in practice

Homework 7-2

- Textbook Exercises 9.5