

The Application Layer

File Transfer/FTP, Web/HTTP and Email

School of Software Engineering
South China University of Technology

Dr. Chunhua Chen

chunhuachen@scut.edu.cn

2019 Spring

Review

Architecture

- Client-Server model
- Peer-to-Peer (P2P) model
- Hybrid model

Services

- Reliability (data loss)
- Throughput
- Timing (delay)

TCP and UDP

Requirements of selected network applications

Application	Data Loss	Throughput	Time-Sensitive
File transfer/download	No loss	Elastic	No
Email	No loss	Elastic	No
Web documents	No loss	Elastic (few kbps)	No
Internet telephony/ Video conferencing	Loss-tolerant	Audio: few kbps–1Mbps Video: 10 kbps–5 Mbps	Yes: 100s of msec
Streaming stored audio /video	Loss-tolerant	Same as above	Yes: few seconds
Interactive games	Loss-tolerant	Few kbps–10 kbps	Yes: 100s of msec
Instant messaging	No loss	Elastic	Yes and no

APP protocol

- Types of messages
- Syntax
- Semantics
- actions

```
GET /img/bd_logo1.png
HTTP/1.1
```

```
Host: www.baidu.com
```

```
Connection: close
```

```
User-agent: Mozilla/5.0
```

```
Accept-language: zh-CN
```

```
HTTP/1.1 200 OK
```

```
Accept-Ranges: bytes
```

```
Age: 3858
```

```
Cache-Control: max-age=315360000
```

```
Content-Length: 7877
```

```
Content-Type: image/png
```

```
Date: Wed, 30 Mar 2016 02:41:35 GMT
```

```
<html>
```

```
<head>
```

```
<title>bd_logo1.png (540×258)</title>
```

```
</head>
```

```
<body style="margin: 0px; background: #0e0e0e;">
```

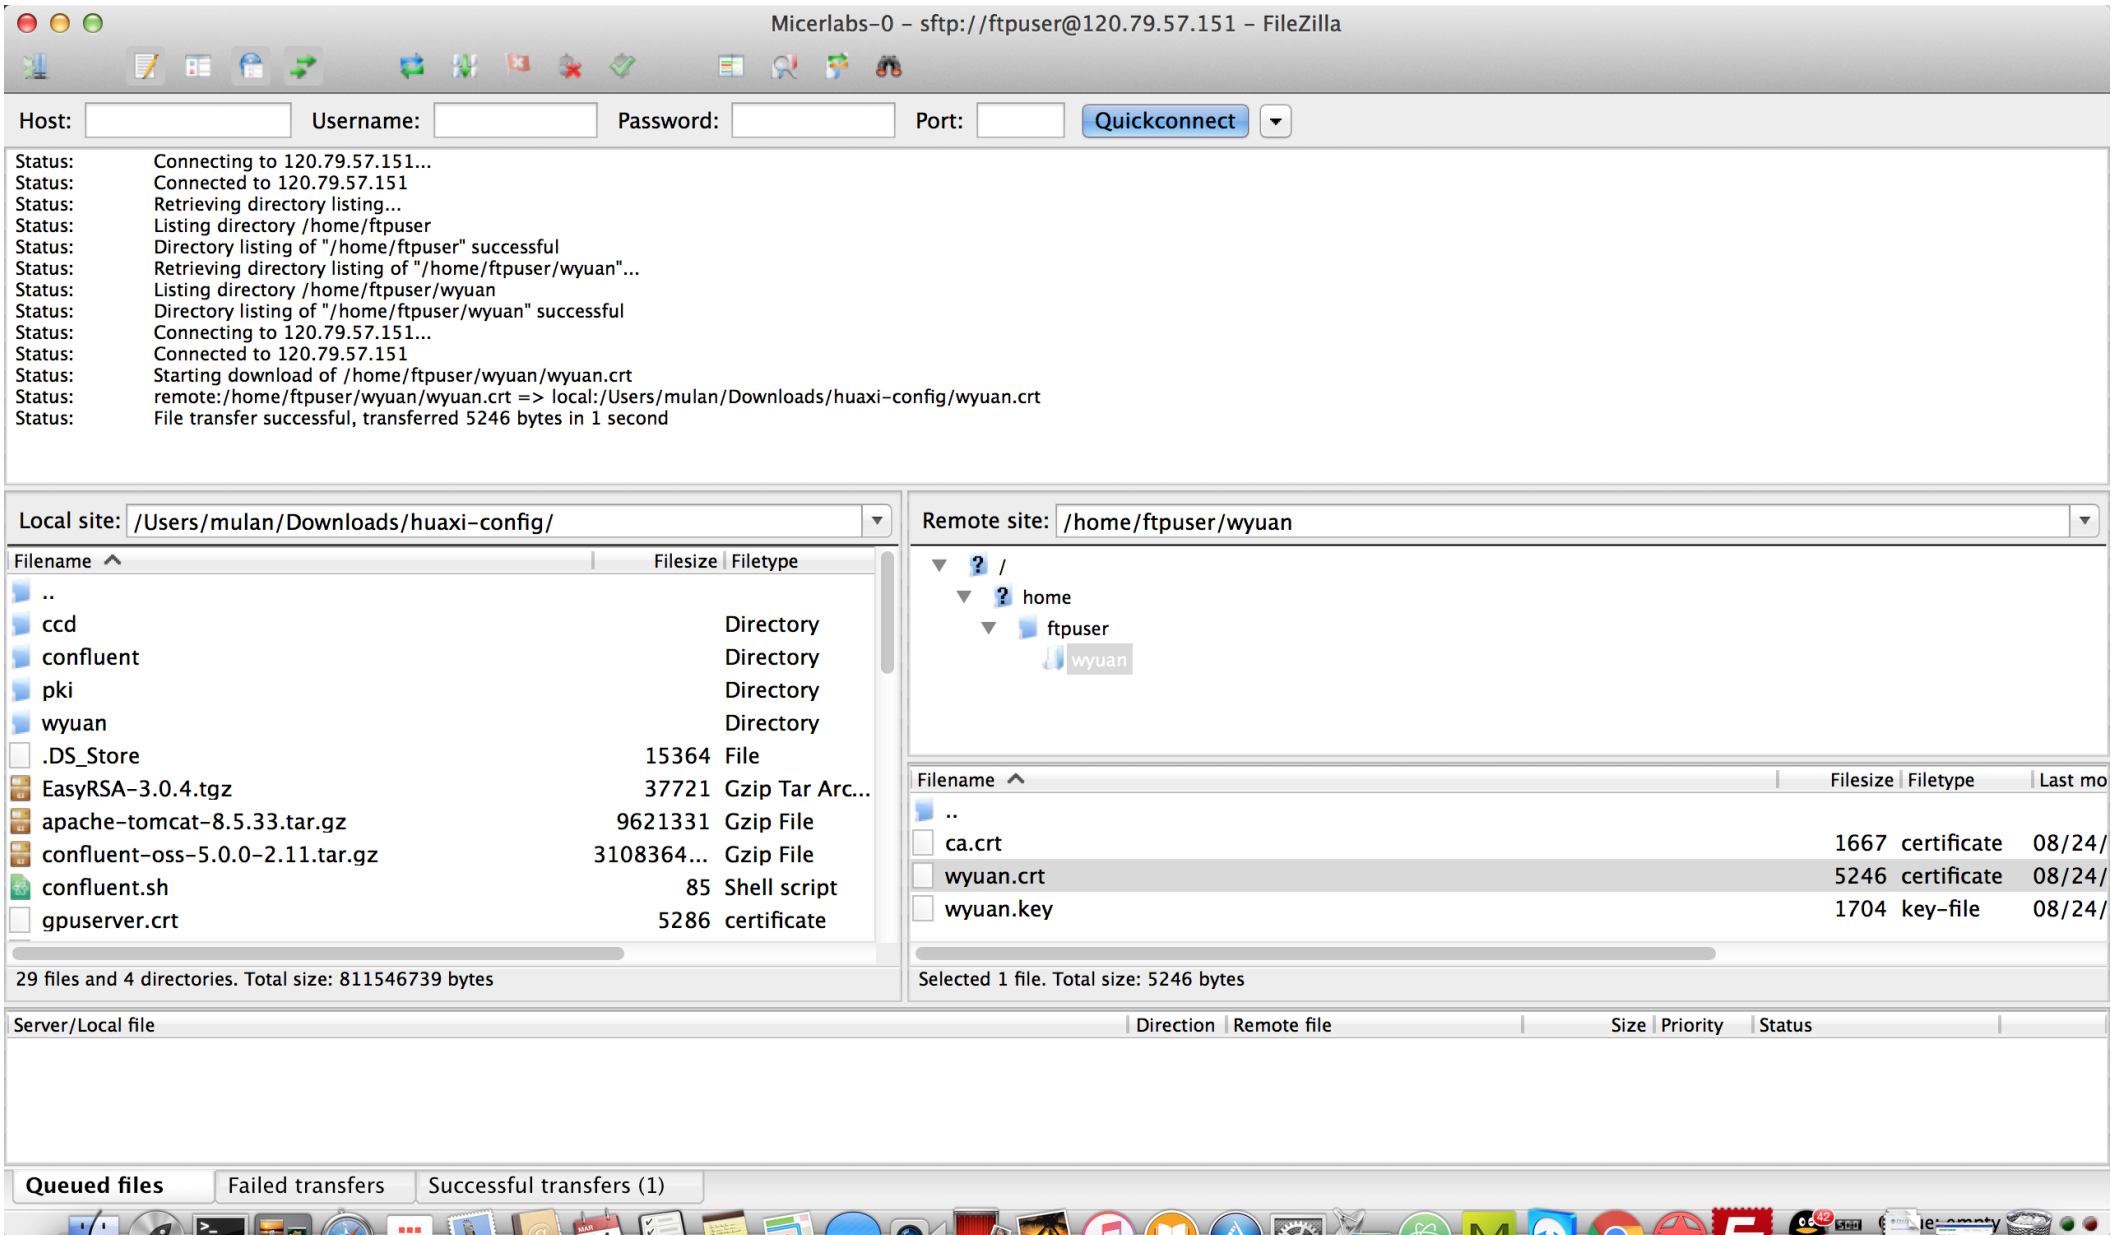
```

```

```
</body>
```

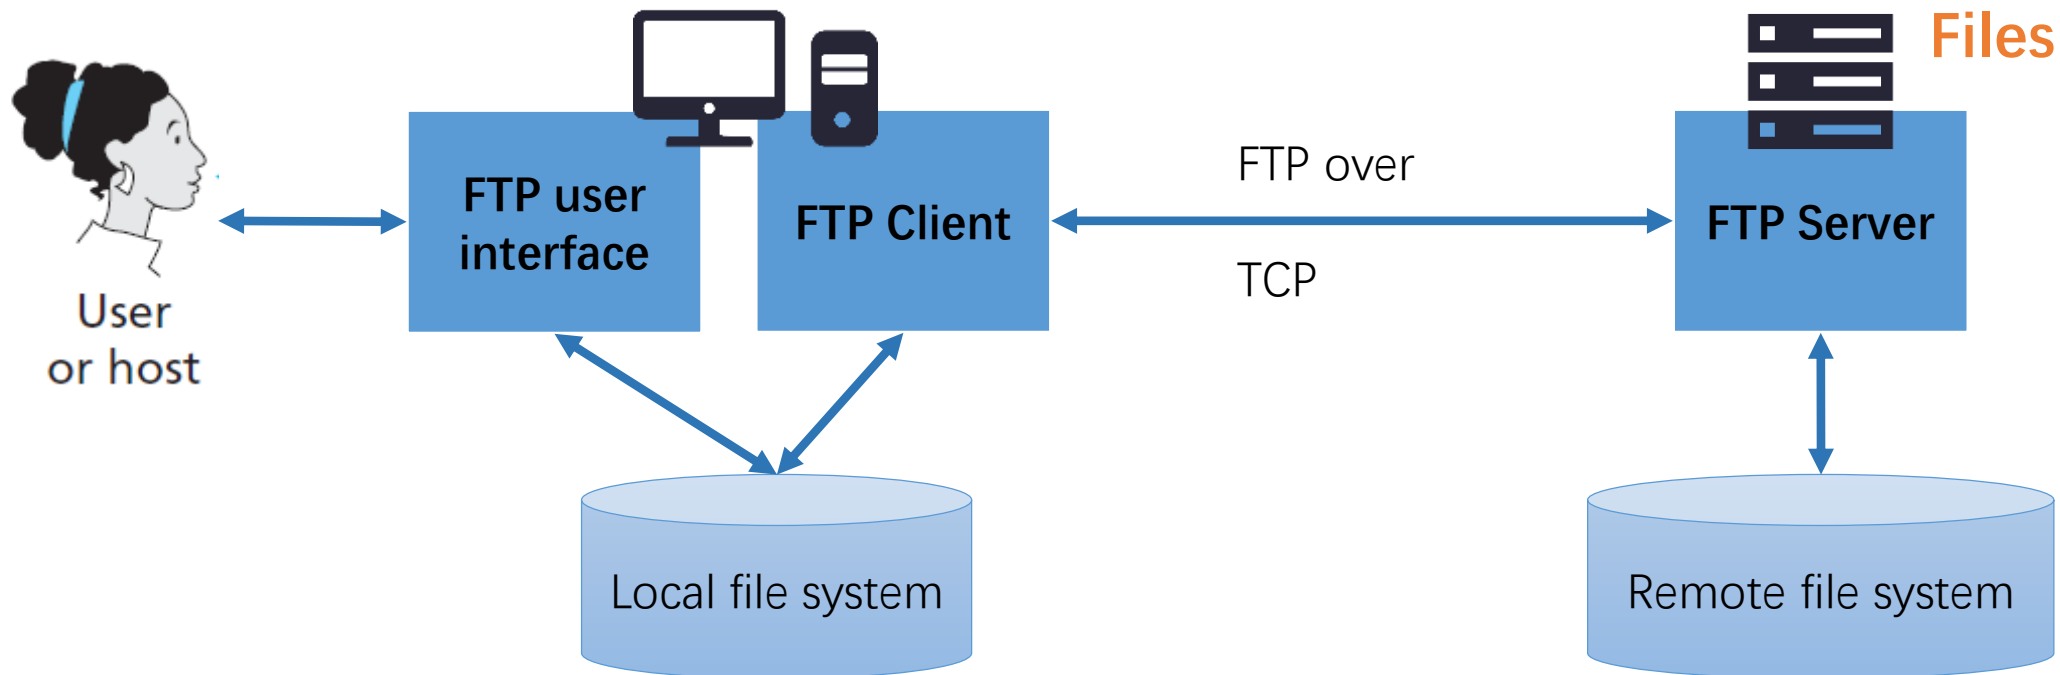
```
</html>
```

How to download a file from a remote server?




File Transfer: FTP

- Client-Server model
- Run on top of TCP, used for download and upload files to FTP Server



FTP: Stateful control and data connection




- The commands, from client to server, and replies, from server to client, are sent across the control connection in 7-bit ASCII format.
- Each command consists of four uppercase ASCII characters, some with optional arguments.
- **USER username**: Used to send the user identification to the server.
- **PASS password**: Used to send the user password to the server.
- **LIST**: Used to ask the server to send back a list of all the files in the current remote directory.
- **RETR filename**: Used to retrieve (that is, get) a file from the current directory of the remote host. This
- **STOR filename**: Used to store (that is, put) a file into the current directory
- of the remote host.

The history of Web

- Invented by Sir Tim Berners-Lee, 1989

How big is the Internet?


@<http://www.internetlivestats.com/total-number-of-websites/>



Sir Tim Berners-Lee

My Visions of Web

- URL (Uniform Resource Locator): Assign an unique address for every resources on the Internet
- HTML: structure relevant resources in to a single document (known as HTML file or Web page) using Hyper Text Mark-up Language (HTML), inform of links (display names associated with URLs)
- HTTP: use Hyper Text Transfer Protocol (HTTP) to download resources on demand.
- A resource also called an object, can be an HTML file, a JPEG image, a Java applet, or a video clip, etc.
- Two components:
 - Web servers, storing resources with addressable URLs
 - Web browsers used by end users to download and use resources from web servers



The URL (or URI)

Addressing every resources in the Web

http://www.baidu.com/img/bd_logo1.png

Host Name

globally unique

Path Name

Unique within Host

What about relationship between a host name and IP?

- The Domain Name System

The Web Documents (or Pages)

<html>

<body>

<p>

[百度图标](http://www.baidu.com/img/bd_logo1.png)

是一个指向百度图标的链接。</p>

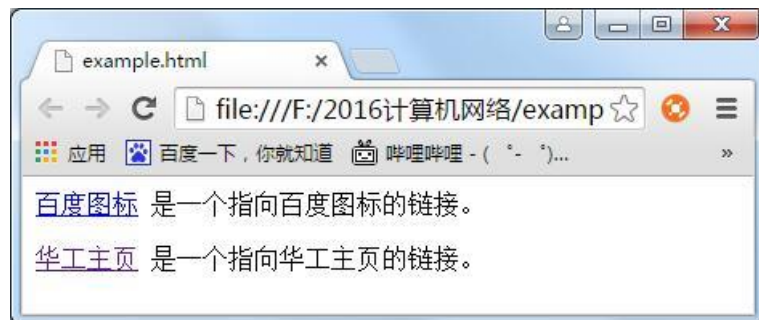
<p>

[华工主页](http://www.scut.edu.cn/)

是一个指向华工主页的链接。</p>


</body>

</html>



The HTTP Protocol

message types, message syntax and message semantics



HTTP 1.0: RFC 1945
HTTP 1.1: RFC 2068

GET /img/bd_logo1.png HTTP/1.1

Host: www.baidu.com

Connection: close

User-agent: Mozilla/5.0

Accept-language: zh-CN

request

HTTP/1.1 200 OK

Accept-Ranges: bytes

Age: 3858

Cache-Control: max-age=315360000

Content-Length: 7877

Content-Type: image/png

Date: Wed, 30 Mar 2016 02:41:35 GMT

ETag: "1ec5-502264e2ae4c0"

Expires: Sat, 28 Mar 2026 02:41:35 GMT

Last-Modified: Wed, 03 Sep 2014 10:00:27 GMT

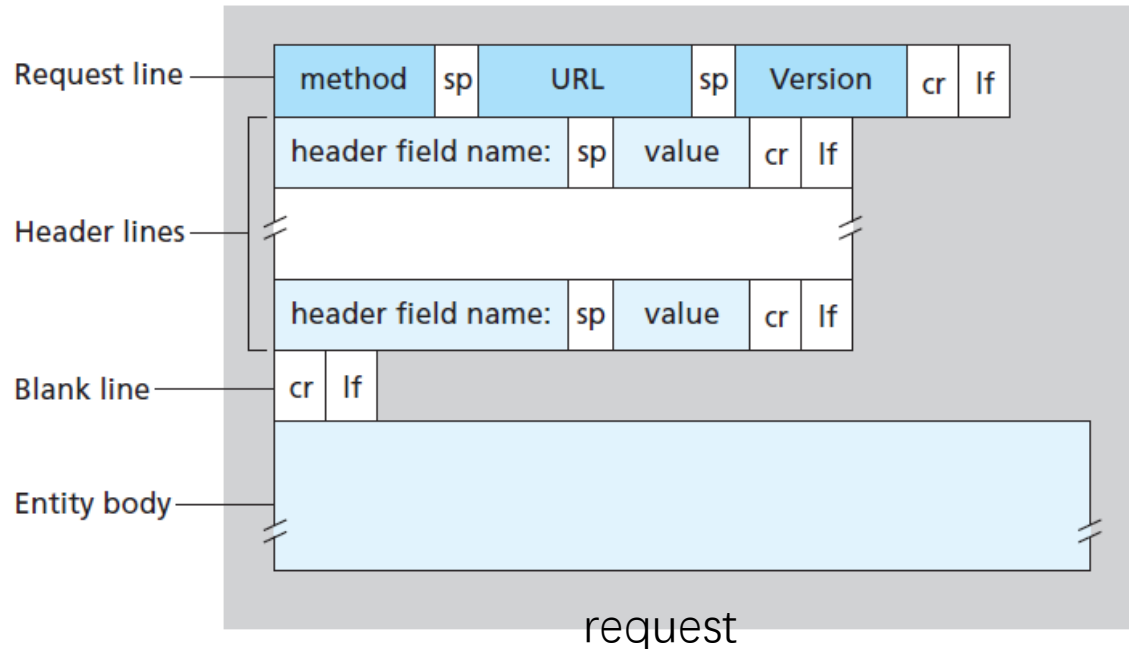
Server: Apache

X-Cache-Lookup: HIT from jpvip175.nydus-vpn.com:3222

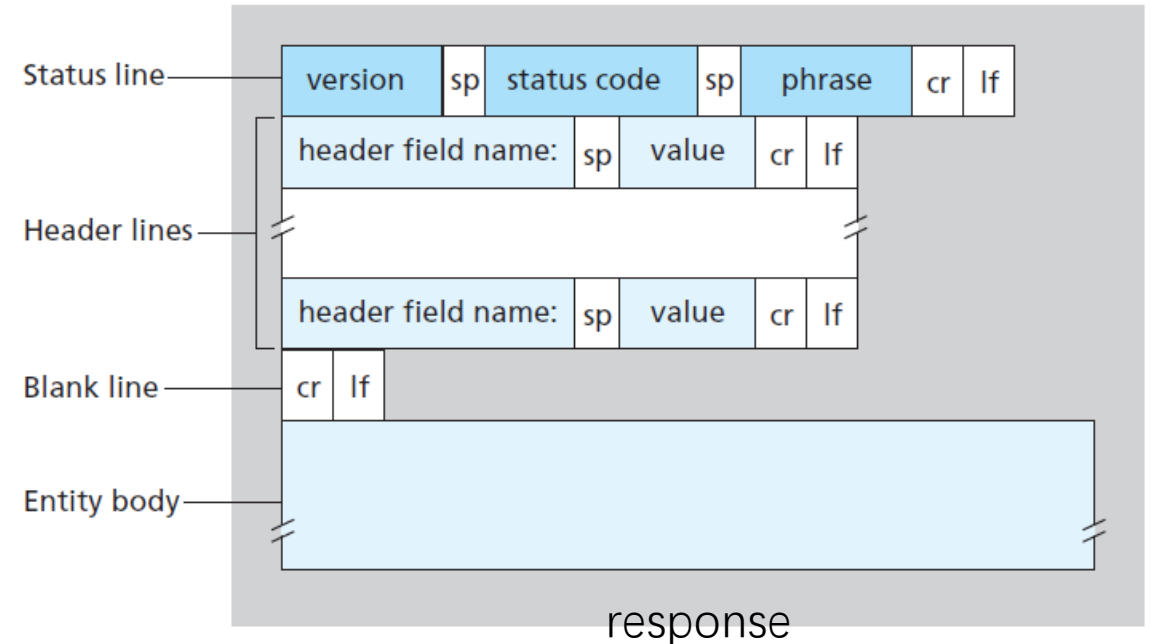
response

The HTTP Protocol

message types, message syntax and message semantics



Method: GET, POST, HEAD, PUT, and DELETE



Status code

200 OK

301 Moved Permanently

400 Bad Request


404 Not Found

505 HTTP Version Not Supported

The HTTP Protocol

stateless connections

- Stateless?
 - server maintains no information about past client requests
 - Can the server trace a particular client?
- Protocols that maintain **state** are complex!
 - past history (state) must be maintained
 - if server/client crashes, their views of state may be inconsistent, must be reconciled



Communication 1 and 2 come from the same client, can server know that? If yes, how?

The HTTP Protocol

non-persistent or persistent connections

- Non-persistent: HTTP 1.0
 - each request/response pair be sent over a **separate** TCP connection
- Persistent: HTTP 1.1
 - multiple request/response pairs be sent over the **same** TCP connection

Non-persistent HTTP

try www.scut.edu.cn/index.html

1a. HTTP client initiates TCP connection to HTTP server (process) at `www.someSchool.edu` on port 80

1b. HTTP server at host `www.someSchool.edu` waiting for TCP connection at port 80. "accepts" connection, notifying client

2. HTTP client sends HTTP *request message* (containing URL) into TCP connection socket. Message indicates that client wants object `/index.html`

3. HTTP server receives request message, forms *response message* containing requested object, and sends message into its socket

4. HTTP server requests to close TCP connection.


5. HTTP client receives response message containing html file, displays html. Parsing html file, finds 10 referenced jpeg objects

6. Steps 1-5 repeated for each of 10 jpeg objects

Non-persistent HTTP


response time

- Definition of RTT
 - time to send a small packet to travel from client to server and back.
- Response time
 - one RTT to initiate TCP connection
 - one RTT for HTTP request and first few bytes of HTTP response to return
 - file transmission time
- total = $2RTT + \text{transmit time}$



Persistent HTTP

- Non-persistent HTTP issues:
 - requires 2 RTTs per object
 - OS overhead for each TCP connection
 - browsers often open parallel TCP connections to fetch referenced objects
- Persistent HTTP
 - server leaves connection open after sending response
 - subsequent HTTP messages between same client/server sent over open connection



Persistent **without pipelining**:

- client issues new request only when previous response has been received
- one RTT for each referenced object

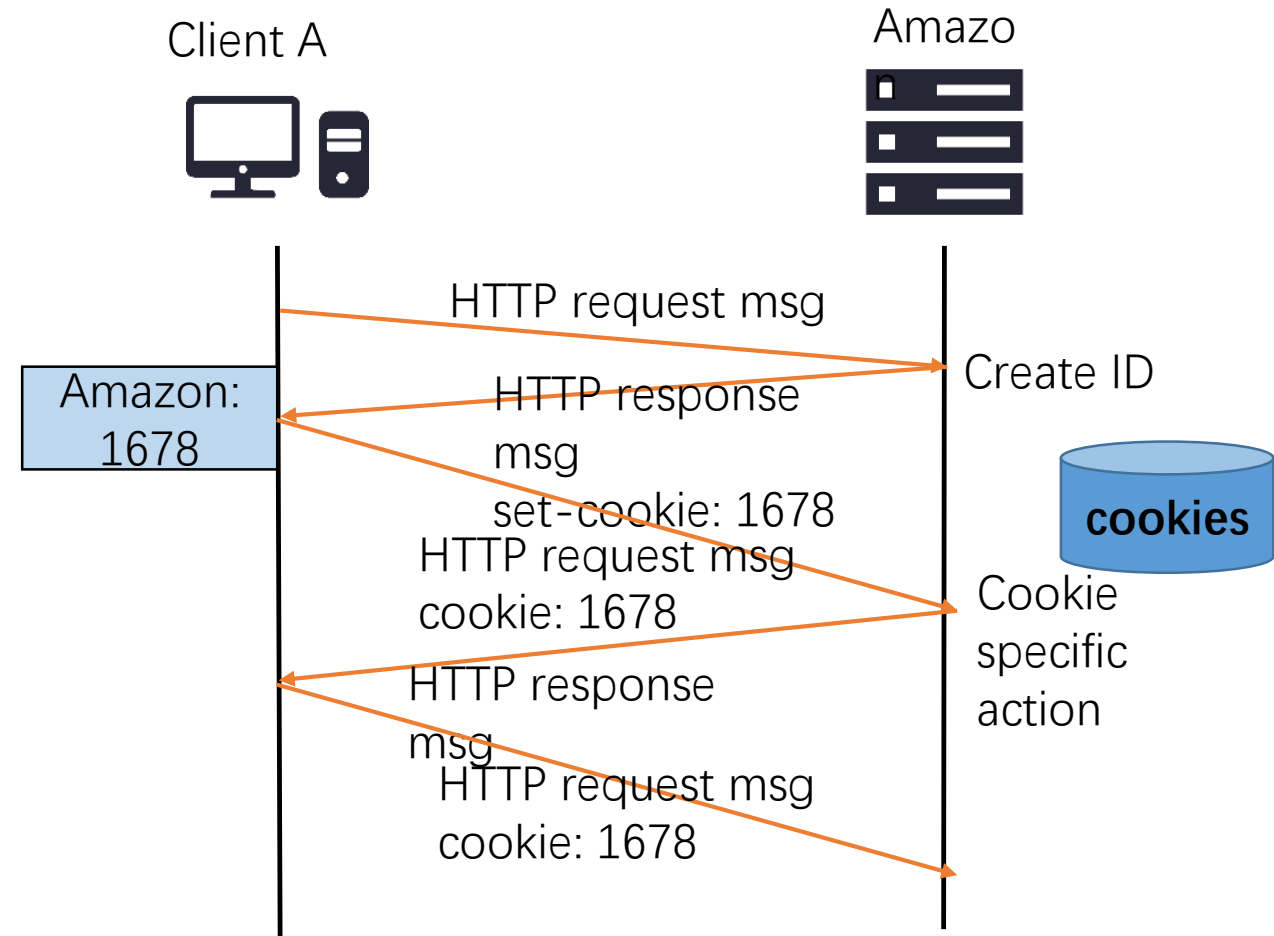
Persistent **with pipelining**:

- default in HTTP/1.1
- client sends requests as soon as it encounters a referenced object
- as little as one RTT for all the referenced objects

User-Server Interaction: Cookies

how can server trace users

- Cookie is a value create and maintained by a server, sent to and maintained by a client.
- In every sequent requests, the client sends cookie with the requests for the server to identify the client.
- Can set life time for a cookie value.



Cookies (continued)

What cookies can bring:

- authorization
- shopping carts
- recommendations
- user session state (Web e-mail)

Cookies and privacy:

- cookies permit sites to learn a lot about you
- you may supply name and e-mail to sites

How to keep “state”:


- protocol endpoints: maintain state at sender/receiver over multiple transactions
- cookies: http messages carry state

Web Caching (Proxy server)

storing content closer to end users

Goal: satisfy client request without involving origin server

- User sets browser: Web accesses via cache
- Browser sends all HTTP requests to cache
 - object in cache: cache returns object
 - else cache requests object from origin server, then returns object to client



More about web caching

- cache acts as both client and server
- typically cache is installed by ISP (university, company, residential ISP)

Why Web caching?

- reduce response time for client request
- reduce traffic on an institution's access link.
- Internet dense with caches: enables “poor” content providers to effectively deliver content (but so does P2P file sharing)

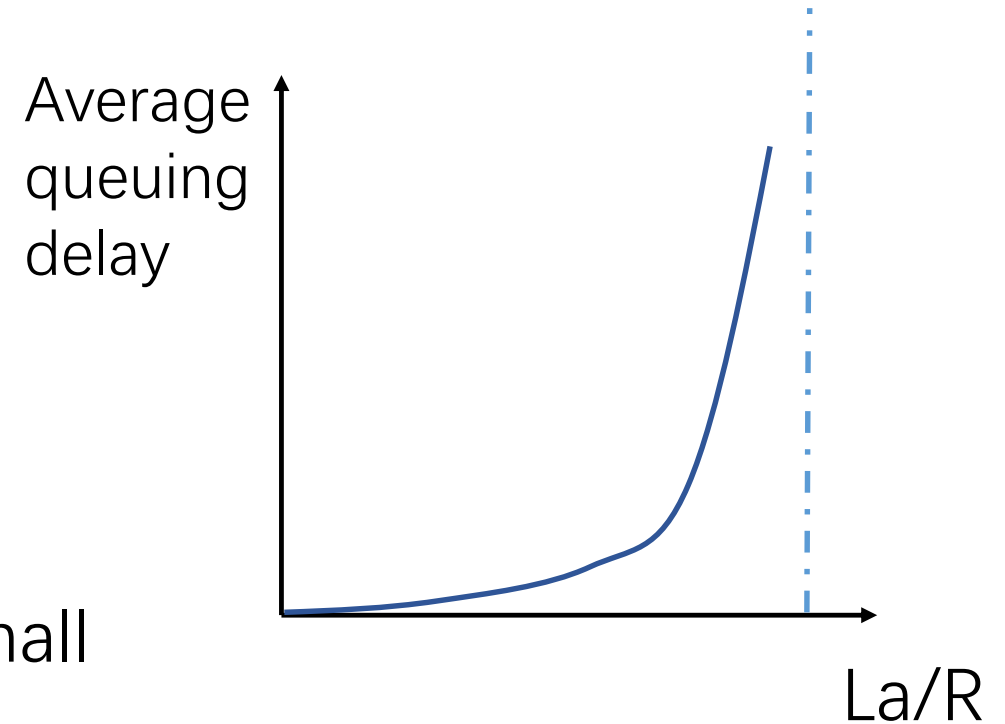
Web cache of computer network vs Storage cache of computer system

Queuing Delay and Traffic Intensity

- R =link bandwidth (bps)
- L =packet length (bits)
- a =average packet arrival rate

Traffic Intensity = La/R

- $La/R \sim 0$: average queuing delay small
- $La/R \rightarrow 1$: delays become large
- $La/R > 1$: more “work” arriving than can be serviced, average delay infinite!



Caching example


the access link is the bottleneck

Assumptions

- average object size = 1Mbits
- avg. request rate from institution's browsers to origin servers = 15/sec
- delay from institutional router to any origin server and back to router = 2 sec

Consequences

- utilization on LAN = 15%
- utilization on access link = 100%
- total delay = Internet delay + access delay + LAN delay
= 2 sec + **minutes** + milliseconds



Caching example (continue)

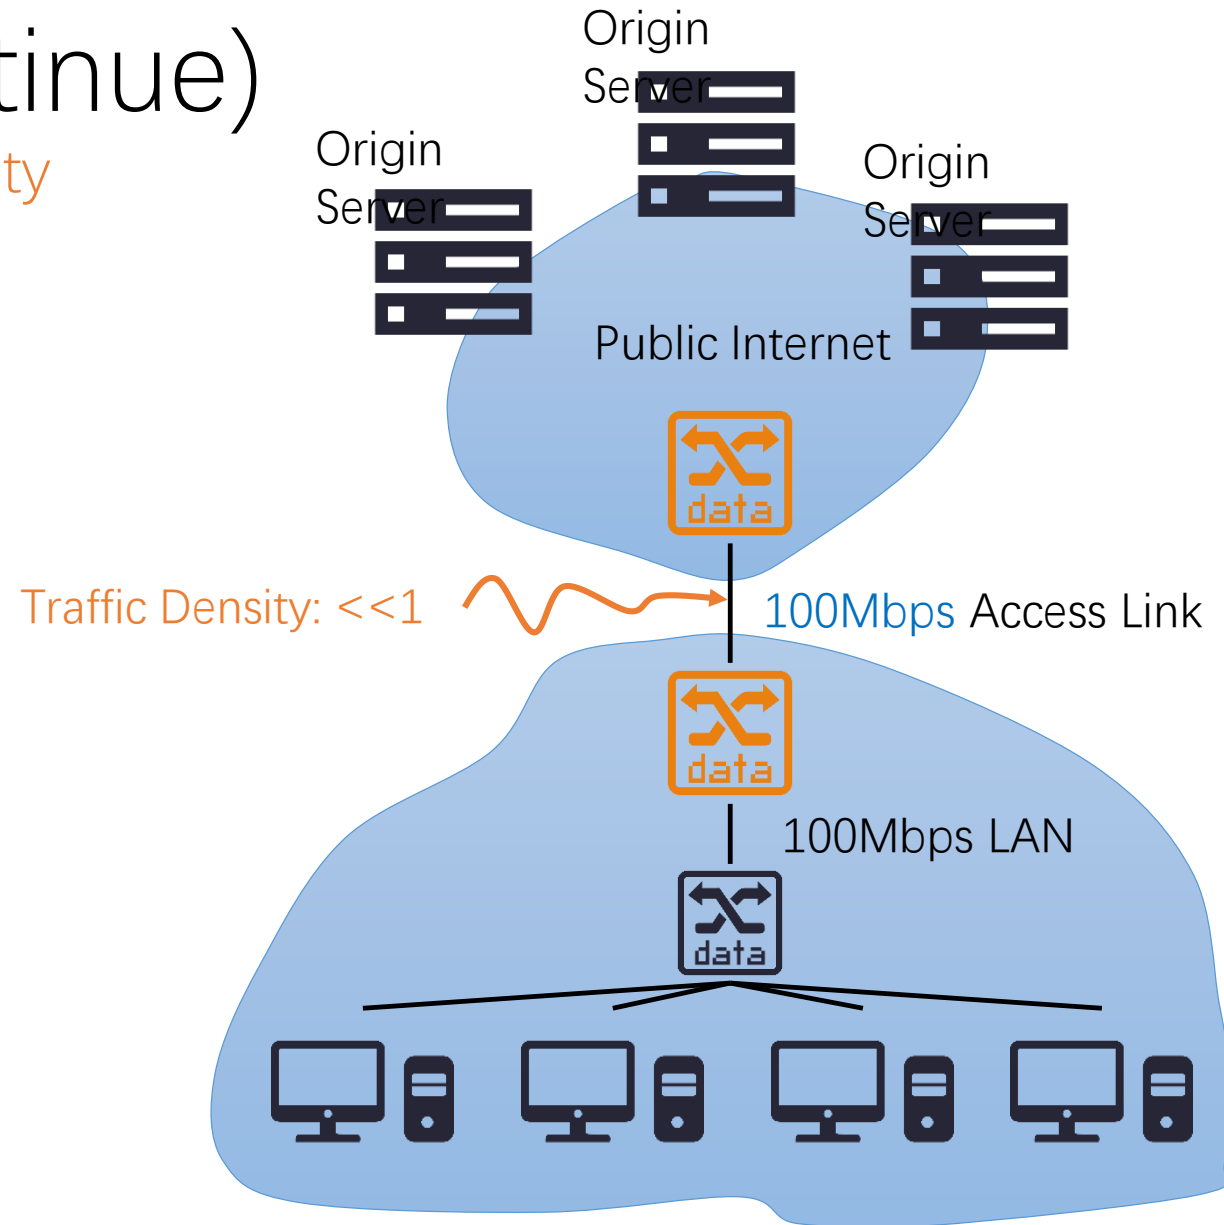
kill the bottleneck by increasing its capability

possible solution

- increase bandwidth of access link to, say, 100 Mbps

consequence

- utilization on LAN = 15%
- utilization on access link = 15%
- Total delay = Internet delay + access delay + LAN delay
= 2 sec + **msecs** + msecs
- often a costly upgrade



Caching example (continue)

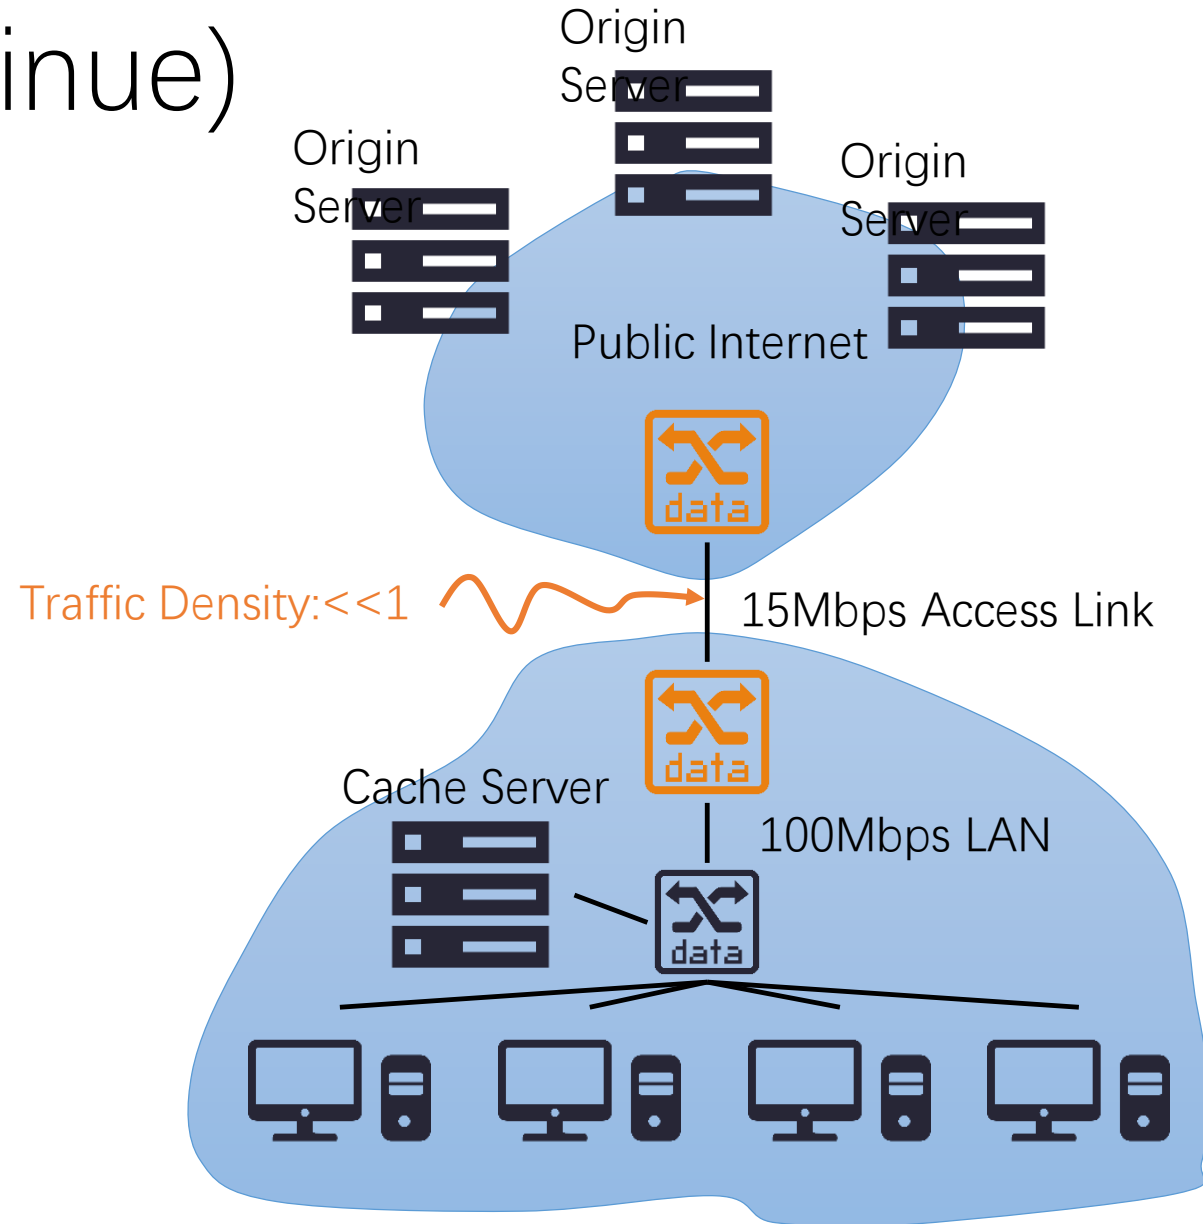
kill the bottleneck by avoiding traffic on it

possible solution: install cache

- suppose hit rate is 0.4

consequence

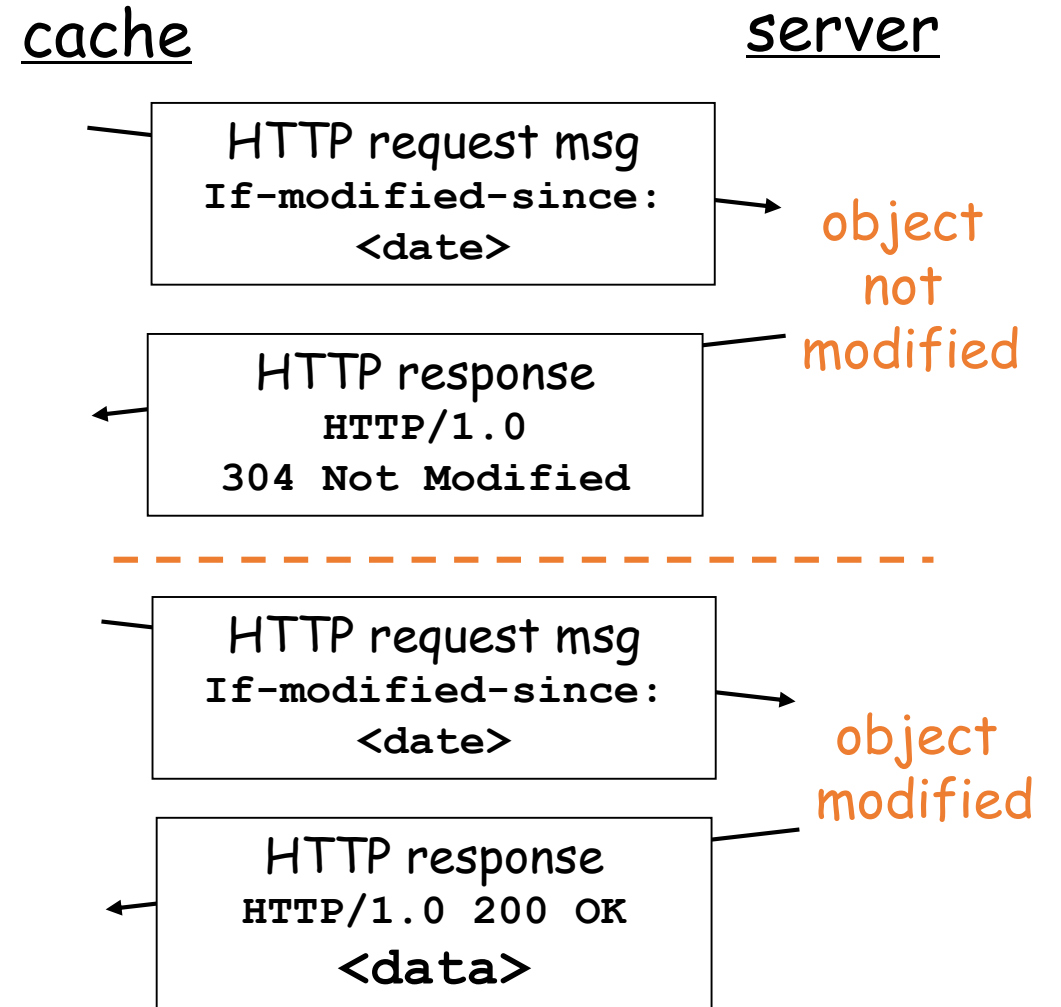
- 40% requests will be satisfied almost immediately
- 60% requests satisfied by origin server
- utilization of access link reduced to 60%, resulting in negligible delays (say 10 msec)
- total avg delay = Internet delay + access delay + LAN delay = $0.6 \times (2) + 0.6 \times 0.01 + 0.4 \times 0.01 \text{ secs} < 1.3 \text{ secs}$



The Conditional GET

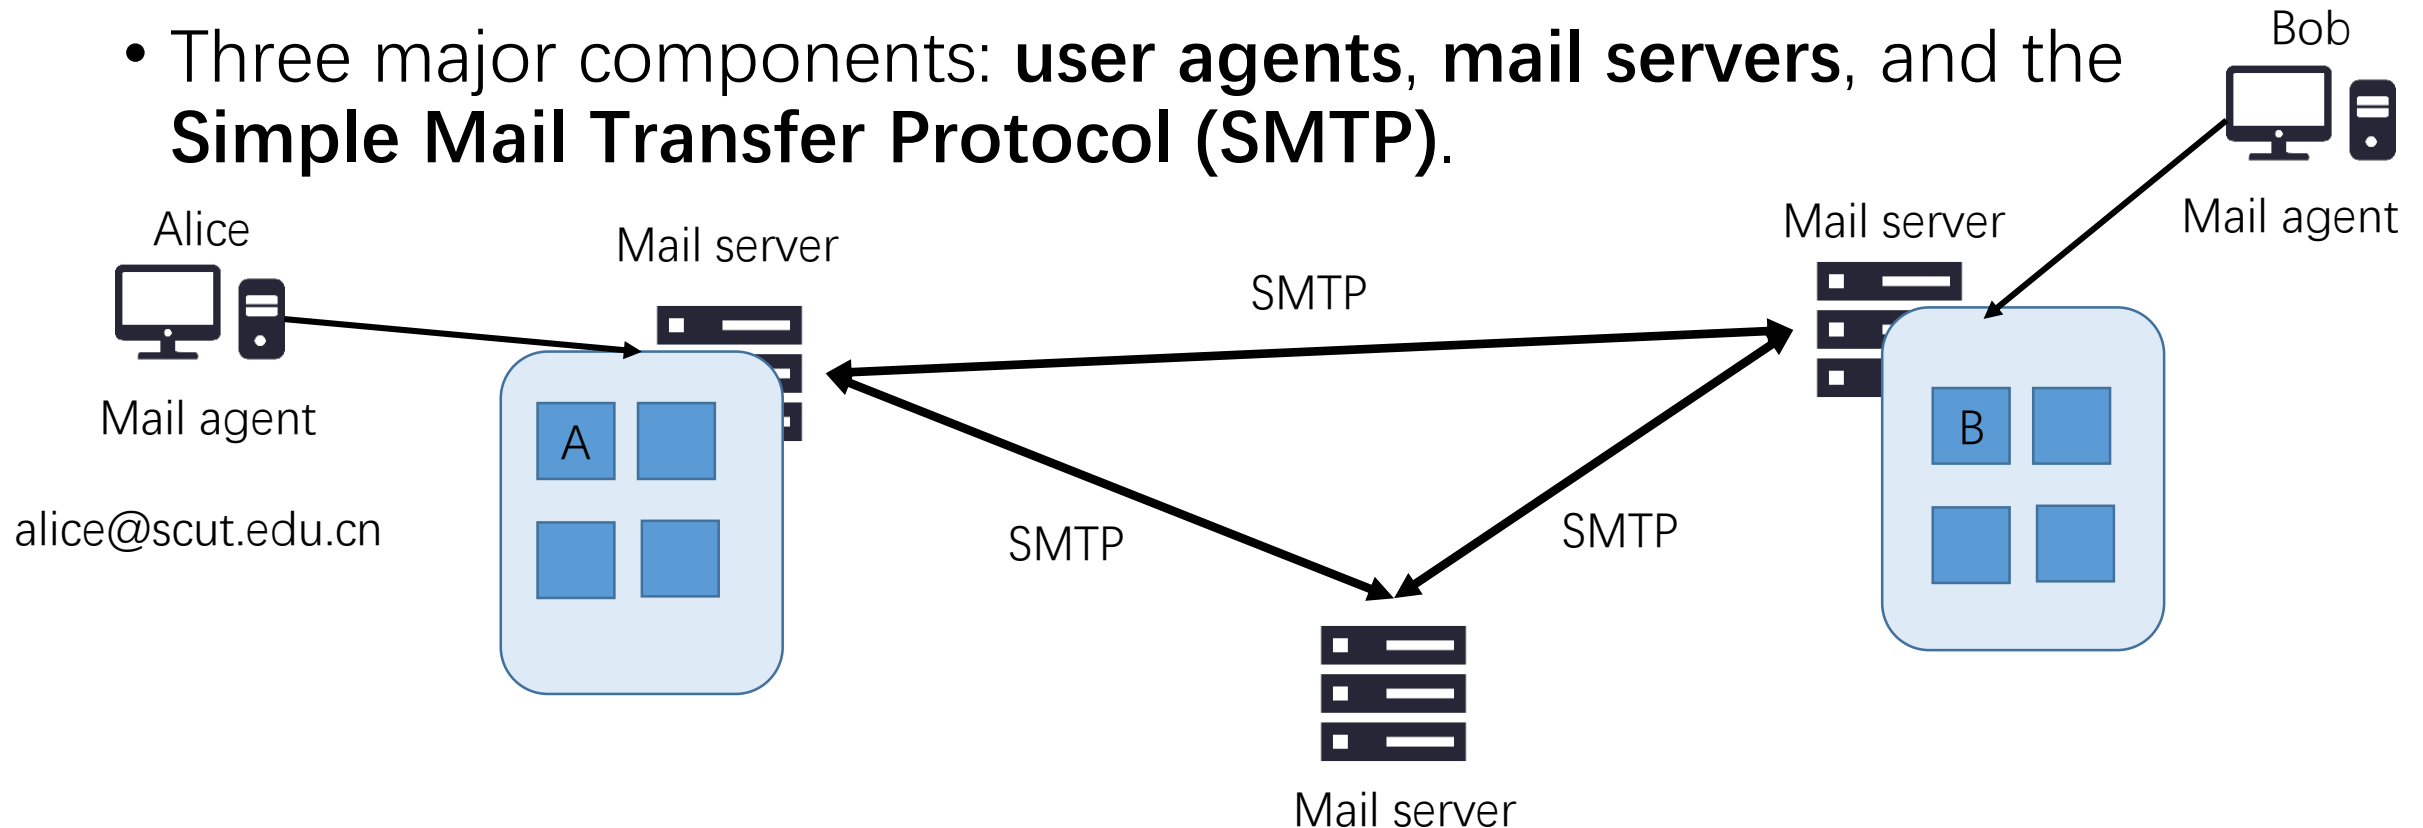
cache at the client

- Goal: don't send object if cache has up-to-date cached version
- cache: specify date of cached copy in HTTP request
`If-modified-since: <date>`
- server: response contains no object if cached copy is up-to-date:
`HTTP/1.0 304 Not Modified`



Electronic Mail in the Internet

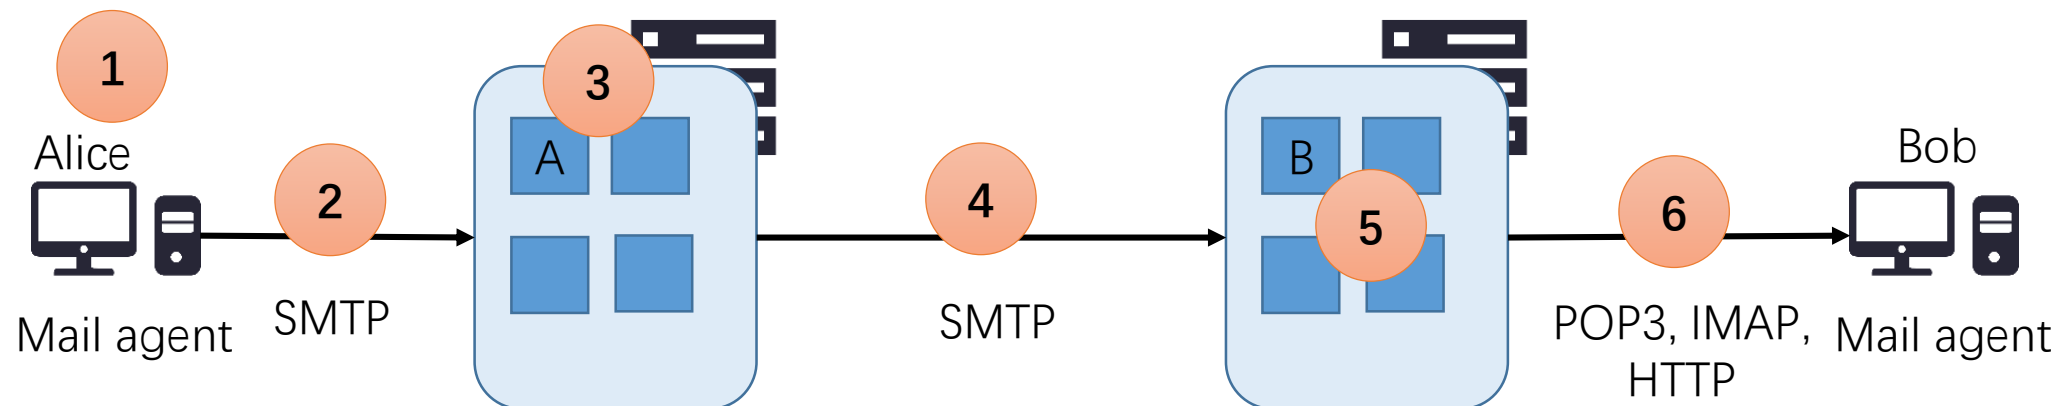
- E-mail is an **asynchronous** communication medium—people send and read messages when it is convenient for them. bob@gmail.com
- Three major components: **user agents**, **mail servers**, and the **Simple Mail Transfer Protocol (SMTP)**.



SMTP @RFC 5321

- Run on TCP, a push protocol model
- HTTP is pull model
- Mail Message Formats @2.4.3
- POP3 and IMAP @ 2.4.4

S: 220 hamburger.edu
C: HELO crepes.fr
S: 250 Hello crepes.fr, pleased to meet you
C: MAIL FROM: <alice@crepes.fr>
S: 250 alice@crepes.fr ... Sender ok
C: RCPT TO: <bob@hamburger.edu>
S: 250 bob@hamburger.edu ... Recipient ok
C: DATA
S: 354 Enter mail, end with "." on a line by itself
C: Do you like ketchup?
C: How about pickles?
C: .
S: 250 Message accepted for delivery
C: QUIT
S: 221 hamburger.edu closing connection



References

- 《图解TCP/IP》与《图解HTTP》(日本人写的)
 - Very good books