

# Chapter 2

## Processes and Threads

2.1 Processes

2.2 Threads

2.3 Inter-process communication

2.4 Classical IPC problems

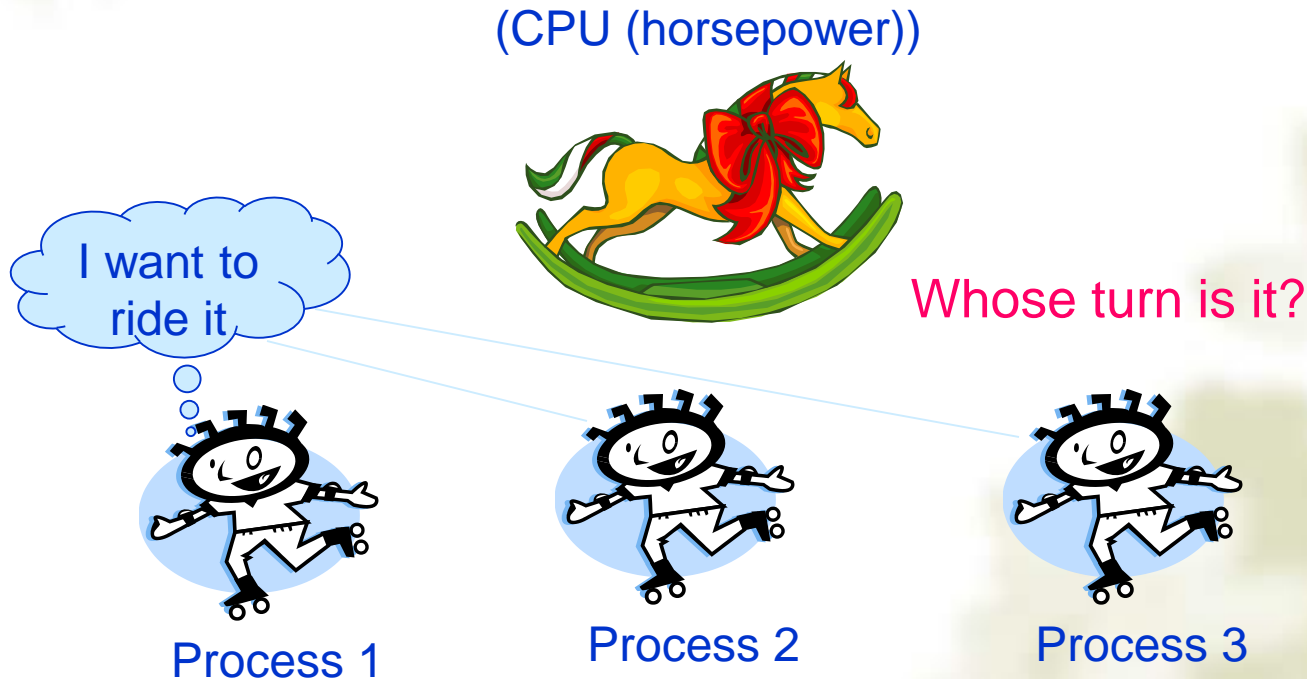
**2.5 Scheduling**

# Content of this lecture

- ❖ What is scheduling?
- ❖ When to Schedule?
- ❖ Basic Scheduling Algorithm
  - ↪ Batch systems
    - ❖ First-Come First-Served
    - ❖ Shortest job first
  - ↪ Interactive systems
    - ❖ Round-robin
    - ❖ Priority scheduling
    - ❖ Multi Queue & Multi-level Feedback
    - ❖ Guaranteed Scheduling, Lottery Scheduling, and Fair Sharing Scheduling
- ❖ Summary

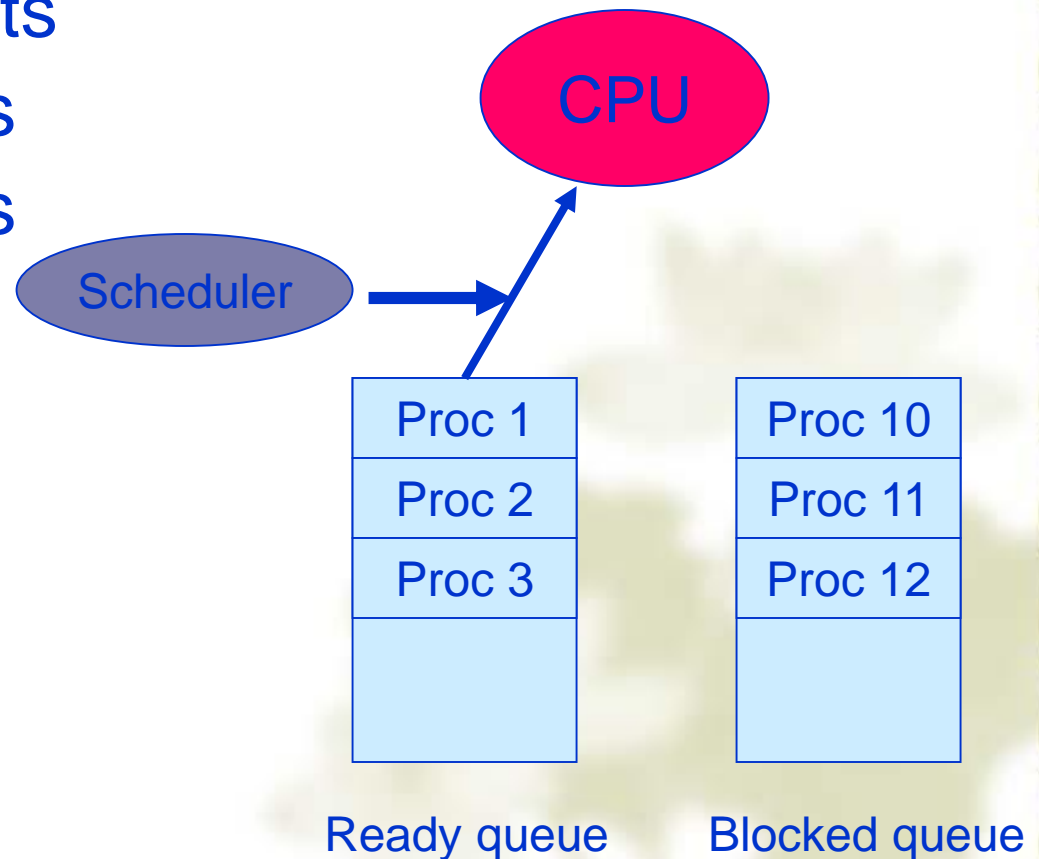
# What is scheduling

- ❖ Deciding which process/thread should occupy the resource (CPU) to run next.



# Process Scheduler

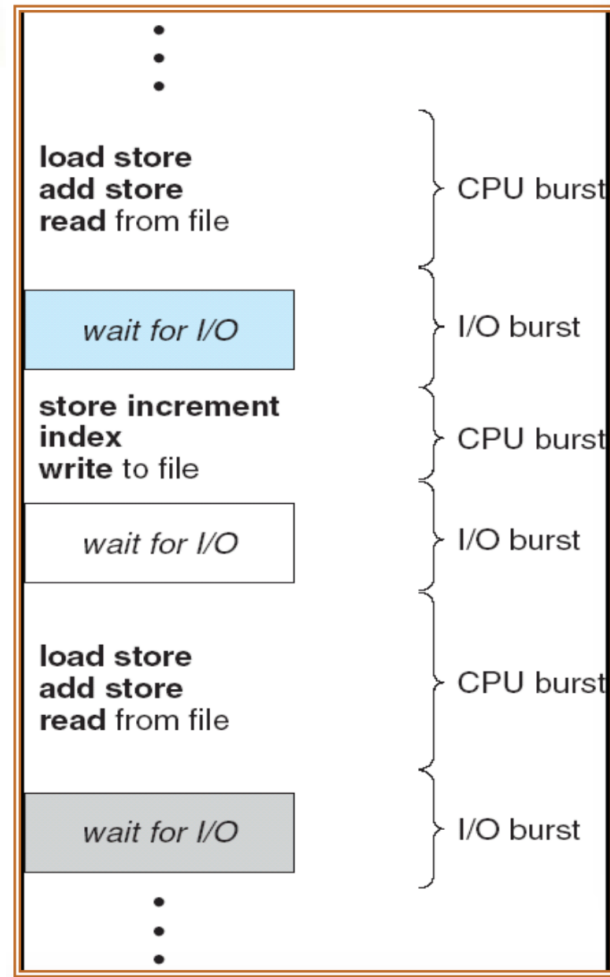
- ❖ Proc1: 20 time units
- ❖ Proc2: 3 time units
- ❖ Proc3: 4 time units
- ❖ Preemptive vs. non-preemptive



# Process Behavior

- ❖ Processes typically consist of
  - ☞ CPU bursts
    - ❖ A period of time when a process needs the CPU is called a CPU burst.
  - ☞ I/O bursts
    - ❖ A period of time when a process needs I/O is called a I/O burst.

# Process Behavior



Alternating Sequence of CPU and I/O Bursts

# Process Behavior

- ❖ Duration and frequency of bursts vary greatly from process to process.

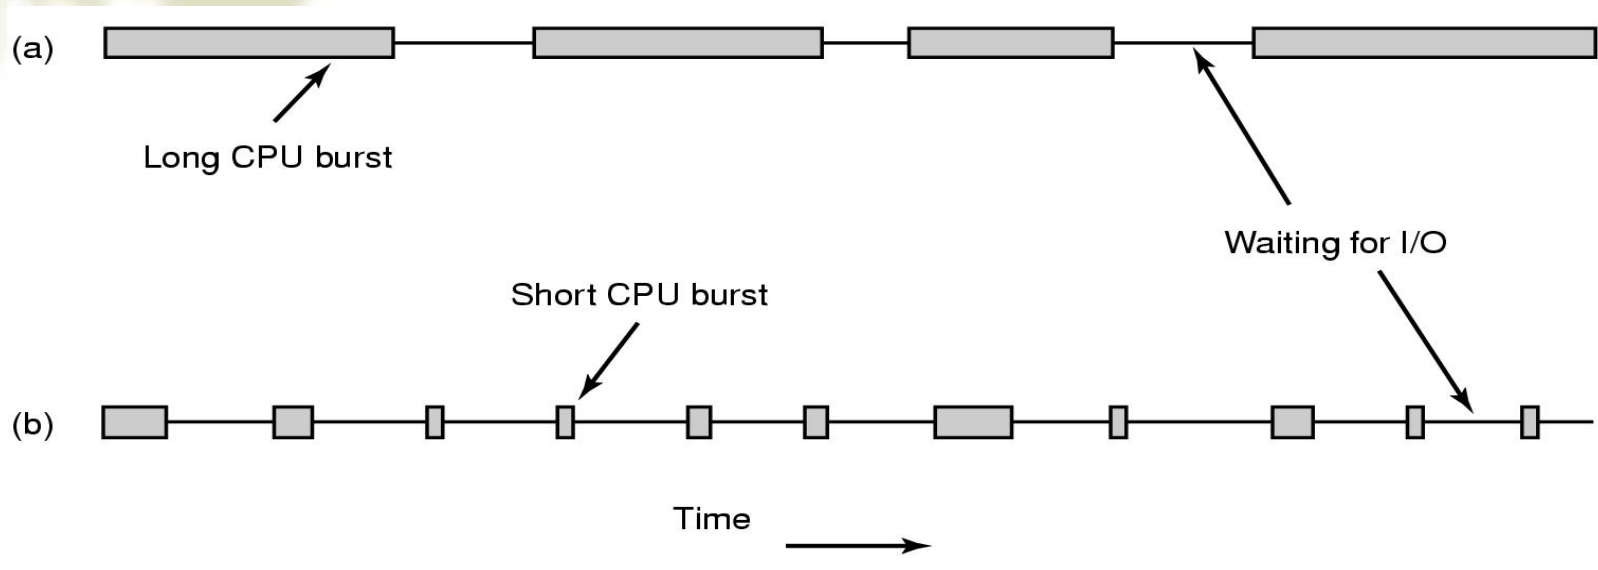
## ☞ CPU-bound Process

- ❖ Long CPU bursts, infrequent I/O waits.
- ❖ E.g. Number crunching tasks, image processing.

## ☞ I/O-bound Process

- ❖ Short CPU bursts, frequent I/O waits.
  - ❖ E.g. A task that processes data from disk, for example, counting the number of lines in a file is likely to be I/O bound.
- ❖ As CPU get faster, processes tend to get more I/O bound.

# Process Behavior



- Bursts of CPU usage alternate with periods of I/O wait
  - a) a CPU-bound process
  - b) an I/O bound process



# CPU Scheduling

## ❖ Problem to solve

- ⌚ When (scheduling opportunity)  
when to allocate CPU to process?
- ⌚ What (scheduling algorithm)  
what is the principle of allocation?
- ⌚ How (context - switch)  
How to allocate CPU? scheduling- process?

# When to schedule?

- ❖ A new process is created
- ❖ The running process exits
- ❖ The running process is blocked
- ❖ I/O interrupt (some processes will be ready)
- ❖ Clock interrupt (every 10 milliseconds)


# Preemptive vs. Non-preemptive

## ❖ Non-preemptive scheduling

⌚ The running process keeps the CPU until it **voluntarily** gives up the CPU

- process exits
- switches to blocked state
- 1 and 4 only (no 3)

⌚ Disadvantage?



## ❖ Preemptive scheduling

⌚ The running process can be interrupted and must release the CPU (can be **forced** to give up CPU)

⌚ Preemptive principles?

# Categories of Scheduling Algorithms

## ❖ Batch System

- ⌚ Non-preemptive algorithms
- ⌚ Preemptive algorithms with long time periods for each process

## ❖ Interactive System

- ⌚ Preemption is essential

## ❖ Real-Time System

- ⌚ Preemption is sometimes not needed

# Scheduling Algorithm

Properties of a GOOD Scheduling Algorithm:

- ❖ Fair (nobody cries)
- ❖ Priority (lady first)
- ❖ Efficiency (make best use of equipment)
- ❖ Encourage good behavior (good boy/girl)
- ❖ Support heavy loads (degrade gracefully)
- ❖ Adapt to different environments (interactive, real-time...)

# Performance Criteria

- ❖ Different Systems, Different Focuses
- ❖ For All Systems

## ↻ Fairness

- ❖ No process should suffer starvation

## ↻ Efficiency

- ❖ Keep resources as busy as possible

## ↻ Policy Enforcement

- ❖ Seeing that stated policy is carried out

# Performance Criteria

## ❖ Batch Systems

### ⌚ Throughput

- ❖ Number of processes that completes in unit time

### ⌚ Turnaround Time (also called elapse time)

- ❖ Amount of time to execute a particular process from the time its entered

### ⌚ Waiting Time

- ❖ amount of time process has been waiting in ready queue

### ⌚ Processor Utilization

- ❖ Percent of time CPU is busy

# Performance Criteria

## ❖ Interactive Systems

### ⌚ Response Time

- ❖ amount of time from when a request was first submitted until first response is produced.

### ⌚ Proportionality

- ❖ meet users' expectation

## ❖ Real-Time Systems

### ⌚ Meeting Deadlines

- ❖ avoid losing data

### ⌚ Predictability

- ❖ Same time/cost regardless of load on the system



# Single Processor Scheduling Algorithms

## ❖ Batch systems

- ↪ First Come First Served (FCFS)
- ↪ Shortest Job First

## ❖ Interactive Systems

- ↪ Round Robin
- ↪ Priority Scheduling
- ↪ Multi Queue & Multi-level Feedback
- ↪ Guaranteed Scheduling
- ↪ Lottery Scheduling
- ↪ Fair Sharing Scheduling


# (1) First Come First Served (FCFS)

- ❖ Also called FIFO. Process that requests the CPU FIRST is allocated the CPU FIRST.
- ❖ Non-preemptive
- ❖ Used in Batch Systems
- ❖ Implementation: FIFO queues
  - ↪ A new process enters the tail of the queue
  - ↪ The scheduler selects from the head of the queue.
- ❖ Performance Metric: **Average Waiting Time.**
- ❖ Given Parameters:
  - ↪ Burst Time (in ms), Arrival Time and Order

# FCFS Example

Process	Burst Time	Order	Arrival Time
P1	24	1	0
P2	3	2	0
P3	4	3	0

The final schedule:



P1 waiting time: 0  
P2 waiting time: 24  
P3 waiting time: 27


The average waiting time:  
 $(0+24+27)/3 = 17$

# FCFS Example (cont.)

Suppose that the processes arrive in the order

$$P_2, P_3, P_1$$

- ❖ The Gantt chart for the schedule is:



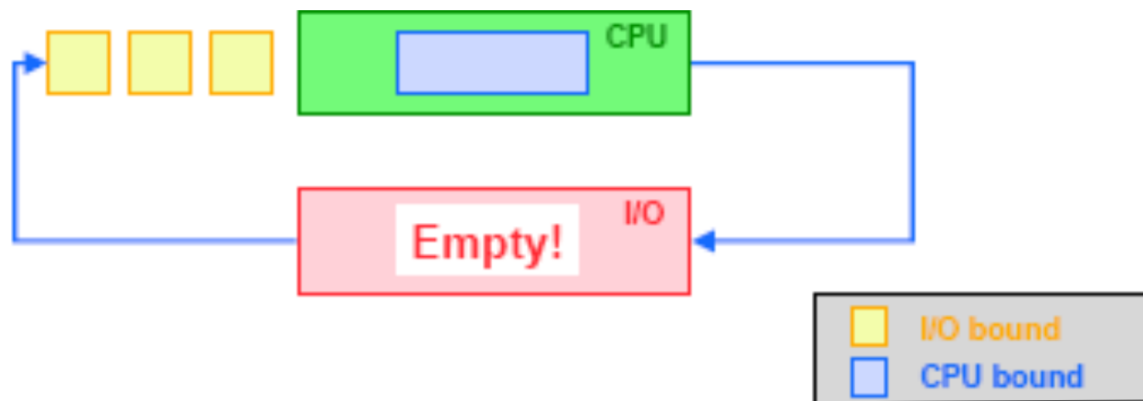
- ❖ Waiting time for  $P_1 = 7$ ;  $P_2 = 0$ ;  $P_3 = 3$
- ❖ Average waiting time:  $(7 + 0 + 3)/3 = 3.3$
- ❖ Much better than previous case.
- ❖ *Convoy effect*: short process behind long process

# Problems with FCFS

- ❖ Non-preemptive
- ❖ Not optimal AWT (Average Waiting Time)
- ❖ Cannot utilize resources in parallel:
  - ⌘ Assume 1 process CPU bounded and many I/O bounded processes
  - ⌘ result: Convoy effect, low CPU and I/O Device utilization
  - ⌘ Why?

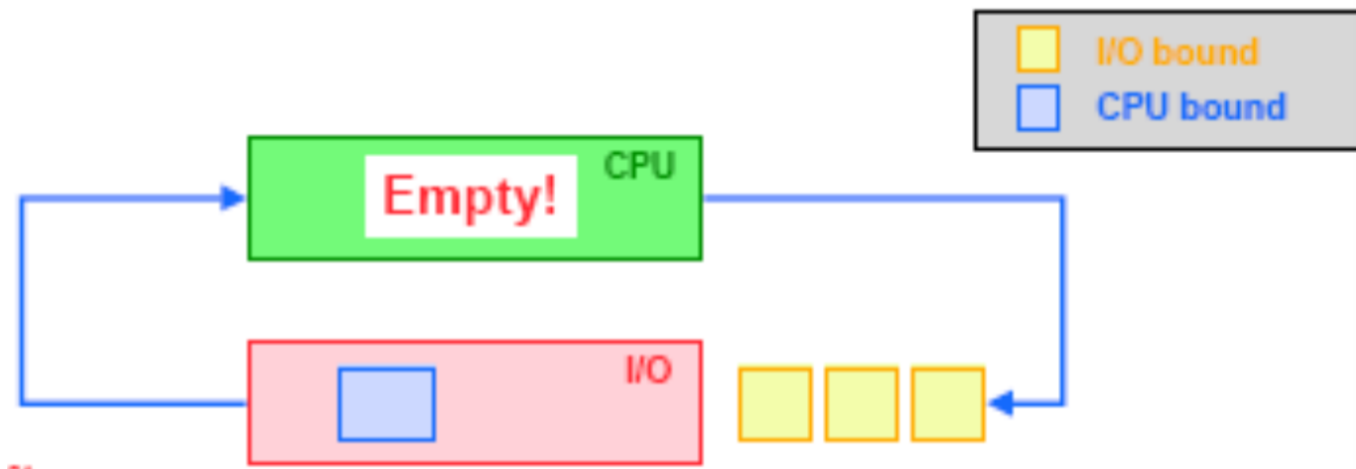
# Why Convoy Effects?

- ❖ Consider  $n-1$  jobs in system that are I/O bound and 1 job that is CPU bound.
- ❖ I/O bound jobs pass quickly through the ready queue and suspend themselves waiting for I/O.
- ❖ CPU bound job arrives at head of queue and executes until complete.
- ❖ I/O bound jobs rejoin ready queue and wait for CPU bound job to complete.
- ❖ I/O devices idle until CPU bound job completes.



# Why Convoy Effects?

- ❖ When CPU bound job complete, other processes rush to wait on I/O again.
- ❖ CPU becomes idle.



## (2) Shortest Job First (SJF)

- ❖ Schedule the job with the shortest elapse time first
- ❖ Scheduling in Batch Systems
- ❖ Two types:
  - ↪ Non-preemptive
  - ↪ Preemptive
- ❖ Requirement: the elapse time needs to know in advance
- ❖ Optimal if all the jobs are available simultaneously (provable)
  - ↪ Gives the best possible AWT (average waiting time)



# Non-preemptive SJF: Example

Process	Burst Time	Order	Arrival Time
P1	6	2	0
P2	8	4	0
P3	7	3	0
P4	3	1	0



P4 waiting time: 0  
P1 waiting time: 3  
P3 waiting time: 9  
P2 waiting time: 16

The total time is: 24  
The average waiting time (AWT):  
 $(0+3+9+16)/4 = 7$

# Comparing to FCFS

Process	Burst Time	Order	Arrival Time
P1	6	1	0
P2	8	2	0
P3	7	3	0
P4	3	4	0



P1 waiting time: 0  
P2 waiting time: 6  
P3 waiting time: 14  
P4 waiting time: 21

The total time is the same.  
The average waiting time (AWT):  
 $(0+6+14+21)/4 = 10.25$   
(comparing to 7)

# SJF is not always optimal

- ❖ Is SJF optimal if all the jobs are not available simultaneously?

Process	Burst Time	Order	Arrival Time
P1	10	1	0
P2	2	2	2




P1 waiting time: 0

P2 waiting time: 8

The average waiting time (AWT):  
 $(0+8)/2 = 4$

# What if the scheduler waits for 2 time units? (Do it yourself)

Process	Burst Time	Order	Arrival Time
P1	10	2	0
P2	2	1	2



P1 waiting time: 4  
P2 waiting time: 0

The average waiting time (AWT):  
 $(0+4)/2 = 2$

However: waste 2 time units of CPU

# Preemptive SJF

- ❖ Also called **Shortest Remaining Time First**
  - ↪ Schedule the job with the shortest remaining time required to complete
- ❖ Requirement: the elapse time needs to be known in advance

# Preemptive SJF: Same Example

Process	Burst Time	Order	Arrival Time
P1	10	1	0
P2	2	2	2



P1 waiting time:  $4 - 2 = 2$

P2 waiting time: 0

The average waiting time (AWT):  
 $(0 + 2) / 2 = 1$

No CPU waste!!!

# A Problem with SJF

## ❖ Starvation

↪ In some condition, a job is waiting for ever

↪ Example: SJF

- ❖ Process A with elapse time of 1 hour arrives at time 0
- ❖ But ever 1 minute from time 0, a short process with elapse time of 2 minutes arrive
- ❖ Result of SJF: A never gets to run

# Interactive Scheduling Algorithms

- ❖ Usually preemptive
  - ↪ Time is sliced into quantum (time intervals)
  - ↪ Scheduling decision is also made at the beginning of each quantum
- ❖ Performance Criteria
  - ↪ Min Response Time
  - ↪ best proportionality
- ❖ Representative algorithms:
  - ↪ Round-robin
  - ↪ Priority-based
  - ↪ Multi Queue & Multi-level Feedback
  - ↪ Guaranteed Scheduling
  - ↪ Lottery Scheduling
  - ↪ Fair Sharing Scheduling




# (3) Round-robin

- ❖ One of the oldest, simplest, most commonly used scheduling algorithm
- ❖ Select process/thread from ready queue in a round-robin fashion (take turns)

Problem:

- Do not consider priority
- Context switch overhead



# Round-robin: Example

Process	Duration	Order	Arrival Time
P1	3	1	0
P2	4	2	0
P3	3	3	0

Suppose time quantum is: 1 unit, P1, P2 & P3 never block

P1 P2 P3 P1 P2 P3 P1 P2 P3 P2



P1 waiting time: 4

P2 waiting time: 6

P3 waiting time: 6

The average waiting time (AWT):  
 $(4+6+6)/3 = 5.33$

# Time Quantum

- ❖ Time slice too large
  - ⌚ FCFS behavior
  - ⌚ Poor response time
- ❖ Time slice too small
  - ⌚ Too many context switches (overheads)
  - ⌚ Inefficient CPU utilization
- ❖ Heuristic: 70-80% of jobs block within time-slice
- ❖ Typical time-slice 10 to 100 ms

## (4) Priority Scheduling


- ❖ Each job is assigned a priority.
- ❖ FCFS within each priority level.
- ❖ Select highest priority job over lower ones.
- ❖ Rational: higher priority jobs are more mission-critical
  - ⌚ Example: display a video film in real time vs. send email
- ❖ Problems:
  - ⌚ May not give the best AWT
  - ⌚ indefinite blocking or starving a process

# Set Priority

- ❖ Two approaches
  - ⌚ Static (for system with well known and regular application behaviors)
  - ⌚ Dynamic (otherwise)
- ❖ Priority may be based on:
  - ⌚ Cost to user
  - ⌚ Importance of user
  - ⌚ Process type
  - ⌚ Requirement to resource
  - ⌚ Aging
  - ⌚ Percentage of CPU time used in last X hours.

# Priority Scheduling: Example

Process	Duration	Priority	Arrival Time
P1	6	4	0
P2	8	1	0
P3	7	3	0
P4	3	2	0



P2 waiting time: 0  
P4 waiting time: 8  
P3 waiting time: 11  
P1 waiting time: 18

The average waiting time (AWT):  
 $(0+8+11+18)/4 = 9.25$   
(worse than SJF:7)

# Priority in Unix

```
yyzhou|csil-linux1|~|[73]% ps -l
```

F	S	UID	PID	PPID	C	PRI	NI	ADDR	SZ	WCHAN	TTY	TIME	CMD
100	S	14828	2047	2045	0	79	4	-	822	rt_sig	pts/1	00:00:00	cs
000	R	14828	23001	2047	0	80	4	-	791	-	pts/1	00:00:00	ps

```
yyzhou|csil-linux1|~|[74]%
```

```
yyzhou|csil-linux1|~|[74]%
```

# Nobody wants to

# Be “nice” in Unix

NICE(1)

FSF

NICE(1)

## NAME

`nice` - run a program with modified scheduling priority

## SYNOPSIS

`nice` [OPTION] [COMMAND] [ARG]...

## DESCRIPTION

Run `COMMAND` with an adjusted scheduling priority. With no `COMMAND`, print the current scheduling priority. `ADJUST` is 10 by default. Range goes from -20 (highest priority) to 19 (lowest).

### -ADJUST

increment priority by `ADJUST` first

### -n, --adjustment=ADJUST

same as `-ADJUST`

`--help` display this help and exit

`--version`

lines 1-25



## (5) Multi-Queue Scheduling

- ❖ Hybrid between priority and round-robin
- ❖ Used in scenarios where the processes can be classified into groups based on property like process type, CPU time, IO access, memory size, etc.
- ❖ Split the Ready Queue in several queues, processes assigned to one queue **permanently**
- ❖ Each queue with its own scheduling algorithm
  - ✿ E.g. interactive processes: RR, background processes: FCFS/SRTF

# Multi-Queue Scheduling

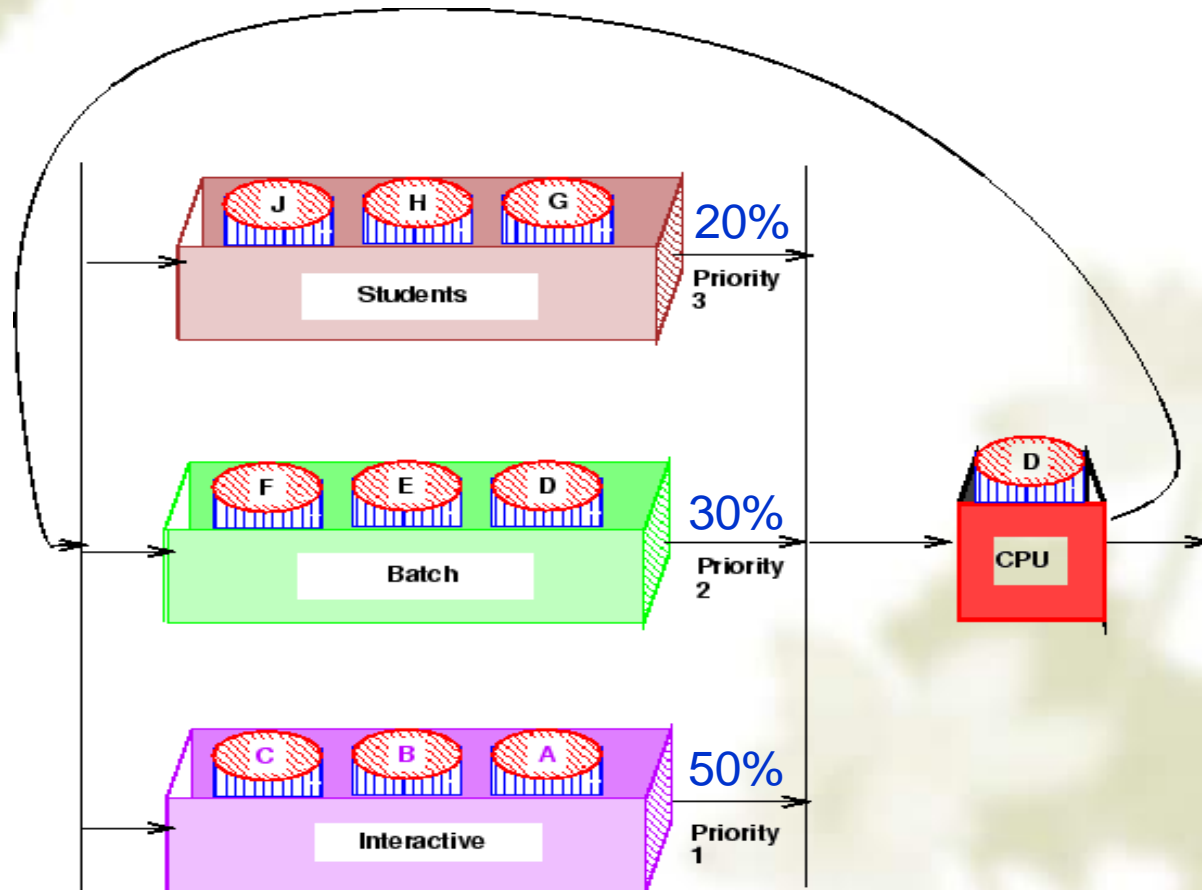
## ❖ Example

- ↪ System processes (highest priority)
- ↪ Interactive programs (Round Robin)
- ↪ Background Processes (FCFS)
- ↪ Student Processes

## ❖ Scheduling between queues

- ↪ Fixed Priorities (Possibility of starvation)
- ↪ % CPU spent on queue

# Multi-Queue Scheduling: Example



# Real Life Analogy

- ❖ Tasks (to-do list) for poor *Bob*
  - ⌚ Class 1 priority (highest): tasks given by his boss
    - ❖ Finish the project (50%)
  - ⌚ Class 2 priority: tasks for his wife
    - ❖ Buy a valentine present (30%)
  - ⌚ Class 3 priority (lowest): Bob's tasks
    - ❖ Watch TV (20%)

## (6) A Variation: Multi-level Feedback Algorithm

- ❖ Multi-Level Queue with priorities
- ❖ Processes **move** between queues
  - ✧ Start each process in a high-priority queue; as it finishes each CPU burst, move it to a lower- priority queue.
- ❖ Each queue represents jobs with similar CPU usage
- ❖ Jobs in a given queue are executed with a given time-slice
- ❖ Rational:
  - ✧ Once an I/O process completes an I/O request, it should have higher CPU priority.

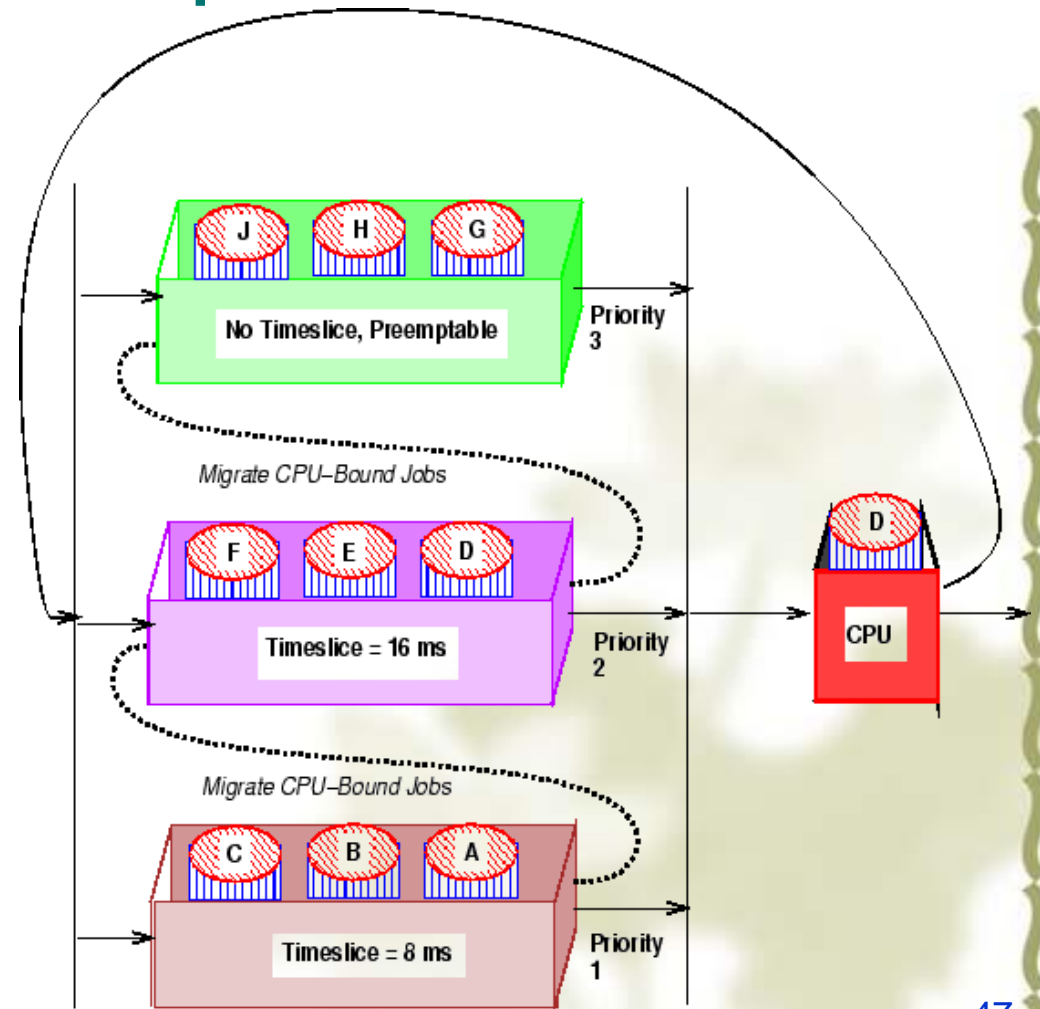
# Multi-level Feedback Algorithm (Details)

- ❖ Example (CTSS):  $Queue_i$  has time-slice  $t \cdot 2^i$ 
  - ↪ If a job in  $Queue_i$  doesn't block by end of time-slice, it is moved to  $Queue_{i+1}$
  - ↪ Lowest priority Queue is FCFS

# Multi-level Feedback Algorithm: Example

## ❖ Three Queues


- ↪  $Q_0$ : time slice 8ms
- ↪  $Q_1$ : time slice 16ms
- ↪  $Q_2$ : FCFS





# Quiz


Show your schedule with timeline and calculate the average turnaround time when use multi-level feedback queue as below (Please take arrival time into account). Note that the priority of the top 2 queues is based on arrival times.



Process ID	Arrival Time	Burst Time
A	0	7
B	2	9
C	5	4
D	7	8
E	8	2



# Solution



Turnaround time:

$$A=18, B=30-2=28, C=23-5=18, D=28-7=21, E=14-8=6$$

Average turnaround time:

$$(18 + 28 + 18 + 21 + 6)/5 = 91/5 = 18.2$$

# Review: Scheduling Algorithms

## ❖ Batch systems

- ↪ First come first served
- ↪ Shortest job first

## ❖ Interactive systems

- ↪ Round-robin
- ↪ Priority scheduling
- ↪ Multi-Queue & Multi-level feedback

# (7) Guaranteed Scheduling (QoS)

- ❖ Make real promises to the users about performance and then live up to them.
- ❖ Example:
  - ↪ with  $n$  processes running, the scheduler makes sure that each one gets  $1/n$  of the CPU cycles.
- ❖ Scheduling:
  - ↪ compute the ratio of actual CPU time consumed to CPU time entitled
  - ↪ Select the one with the lowest ratio

## (8) Lottery Scheduling

- ❖ More commonly used
- ❖ Probability-based
  - ✧ Give processes lottery tickets. At scheduling time, a lottery ticket is chosen at random, and the process holding that ticket gets that resource.
- ❖ Give more tickets for higher priority processes (approximate Priority), or give short jobs more tickets (approximate SJF)
- ❖ Advantages
  - ✧ Simple
  - ✧ Highly responsive
  - ✧ Can support cooperation between processes
  - ✧ Easy to support priority and proportion requirement

# (9) Fair-Share Scheduling

- ❖ Is Round-robin fair?
  - ↪ Yes, it is (from process point of view)
  - ↪ No, it may be not (from user point view)
- ❖ User-based fair share scheduling
  - ↪ Each user gets fair share
- ❖ Example:
  - ↪ Alice has 4 processes: A1, A2, A3, A4
  - ↪ Bob has 1 process: B1
  - ↪ Then A1, A2, A3, A4 are entitled only to 50% CPU, while B1 alone is entitled to 50%
  - ↪ Possible scheduling sequence:  
A1, **B1**, A2, **B1**, A3, **B1**, A4, **B1**, .....

# Scheduling in Real-Time Systems

- ❖ The scheduler makes real promises to the user in terms of deadlines or CPU utilization.
- ❖ Schedulable real-time system

↪ Given

- ❖  $m$  periodic events
- ❖ event  $i$  occurs within period  $P_i$  and requires  $C_i$  seconds


↪ Then the load can only be handled if

$$\sum_{i=1}^m \frac{C_i}{P_i} \leq 1$$

# User-level Thread Scheduling

## Possible Scheduling

- ❖ 50-msec quantum
- ❖ run 5 msec/CPU burst



Possible: A1, A2, A3, A1, A2, A3


Not possible: A1, B1, A2, B2, A3, B3



# Kernel-level Thread Scheduling

## Possible scheduling

- ❖ 50-msec quantum
- ❖ threads run 5 msec/CPU burst



Possible: A1, A2, A3, A1, A2, A3

Also possible: A1, B1, A2, B2, A3, B3



# Summary

- ❖ What is scheduling
- ❖ Scheduling objectives
- ❖ CPU Scheduling
- ❖ FCFS
- ❖ Shortest job first (SJF)
- ❖ Priority
- ❖ Round-robin
- ❖ Multi-Queue
- ❖ Multi-level Feedback
- Scheduling algorithms
  - Guaranteed Scheduling
  - Lottery Scheduling
  - Fair Sharing Scheduling
- Scheduling for
  - Real-time systems
  - Threads

# Homework

- ❖ 45、50
- ❖ Reading assignment: 10.3